

Accessing Databases with JDBC

28



It is a capital mistake to theorize before one has data.

—Arthur Conan Doyle

Now go, write it before them in a table, and note it in a book, that it may be for the time to come for ever and ever.

—The Holy Bible, Isaiah 30:8

Get your facts first, and then you can distort them as much as you please.

—Mark Twain

I like two kinds of men: domestic and foreign.

—Mae West

Objectives

In this chapter you'll learn:

- Relational database concepts.
- To use Structured Query Language (SQL) to retrieve data from and manipulate data in a database.
- To use the JDBC™ API to access databases.
- To use the RowSet interface to manipulate databases.
- To use JDBC 4.0's automatic JDBC driver discovery.
- To use PreparedStatements to create precompiled SQL statements with parameters.
- How transaction processing makes database applications more robust.

Outline



28.1 Introduction	28.6 Instructions for Setting Up a MySQL User Account
28.2 Relational Databases	28.7 Creating Database books in MySQL
28.3 Relational Database Overview: The books Database	28.8 Manipulating Databases with JDBC
28.4 SQL	28.8.1 Connecting to and Querying a Database
28.4.1 Basic SELECT Query	28.8.2 Querying the books Database
28.4.2 WHERE Clause	28.9 RowSet Interface
28.4.3 ORDER BY Clause	28.10 Java DB/Apache Derby
28.4.4 Merging Data from Multiple Tables: INNER JOIN	28.11 PreparedStatements
28.4.5 INSERT Statement	28.12 Stored Procedures
28.4.6 UPDATE Statement	28.13 Transaction Processing
28.4.7 DELETE Statement	28.14 Wrap-Up
28.5 Instructions for Installing MySQL and MySQL Connector/J	28.15 Web Resources

[Summary](#) | [Terminology](#) | [Self-Review Exercise](#) | [Answers to Self-Review Exercise](#) | [Exercises](#)

28.1 Introduction

A **database** is an organized collection of data. There are many different strategies for organizing data to facilitate easy access and manipulation. A **database management system (DBMS)** provides mechanisms for storing, organizing, retrieving and modifying data for many users. Database management systems allow for the access and storage of data without concern for the internal representation of data.

Today's most popular database systems are relational databases, where the data is stored without consideration of its physical structure (Section 28.2). A language called **SQL**—pronounced “sequel,” or as its individual letters—is the international standard language used almost universally with relational databases to perform **queries** (i.e., to request information that satisfies given criteria) and to manipulate data. [Note: As you learn about SQL, you'll see some authors writing “a SQL statement” (which assumes the pronunciation “sequel”) and others writing “an SQL statement” (which assumes that the individual letters are pronounced). In this book we pronounce SQL as “sequel.”]

Some popular **relational database management systems (RDBMSs)** are Microsoft SQL Server, Oracle, Sybase, IBM DB2, Informix, PostgreSQL and MySQL. The JDK now comes with a pure-Java RDBMS called Java DB—Sun's version of Apache Derby. In this chapter, we present examples using MySQL and Java DB.

Java programs communicate with databases and manipulate their data using the **Java Database Connectivity (JDBC™ API)**. A **JDBC driver** enables Java applications to connect to a database in a particular DBMS and allows you to manipulate that database using the JDBC API.



Software Engineering Observation 28.1

Using the JDBC API enables developers to change the underlying DBMS without modifying the Java code that accesses the database.

Most popular database management systems now provide JDBC drivers. There are also many third-party JDBC drivers available. In this chapter, we introduce JDBC and use it to manipulate MySQL and Java DB databases. The techniques demonstrated here can also be used to manipulate other databases that have JDBC drivers. Check your DBMS's documentation to determine whether your DBMS comes with a JDBC driver. If not, third-party vendors provide JDBC drivers for many DBMSs.



Software Engineering Observation 28.2

Most major database vendors provide their own JDBC database drivers, and many third-party vendors provide JDBC drivers as well. For more information visit devapp.sun.com/product/jdbc/drivers.

For more information on JDBC, visit

java.sun.com/javase/technologies/database/index.jsp

which contains JDBC information including the JDBC specification, FAQs, a learning resource center and software downloads to search for JDBC drivers for your DBMS, and

developers.sun.com/product/jdbc/drivers/

which provides a search engine to help you locate drivers appropriate for your DBMS.

28.2 Relational Databases

A **relational database** is a logical representation of data that allows the data to be accessed without consideration of its physical structure. A relational database stores data in **tables**. Figure 28.1 illustrates a sample table that might be used in a personnel system. The table name is **Employee**, and its primary purpose is to store the attributes of an employee. Tables are composed of **rows**, and rows are composed of **columns** in which values are stored. This table consists of six rows. The **Number** column of each row is the table's **primary key**—a column (or group of columns) with a unique value that cannot be duplicated in other rows. This guarantees that each row can be identified by its primary key. Good examples of primary key columns are a social security number, an employee ID number and a part number in an inventory system, as values in each of these columns are guaranteed to be unique. The rows in Fig. 28.1 are displayed in order by primary key. In this case, the rows are listed in increasing order, but we could also use decreasing order.

	Number	Name	Department	Salary	Location
Row	23603	Jones	413	1100	New Jersey
	24568	Kerwin	413	2000	New Jersey
	34589	Larson	642	1800	Los Angeles
	35761	Myers	611	1400	Orlando
	47132	Neumann	413	9000	New Jersey
	78321	Stephens	611	8500	Orlando

Fig. 28.1 | Employee table sample data.

Rows in tables are not guaranteed to be stored in any particular order. As we'll demonstrate in an upcoming example, programs can specify ordering criteria when requesting data from a database.

Each column represents a different data attribute. Rows are normally unique (by primary key) within a table, but particular column values may be duplicated between rows. For example, three different rows in the `Employee` table's `Department` column contain number 413.

Different users of a database are often interested in different data and different relationships among the data. Most users require only subsets of the rows and columns. To obtain these subsets, you use queries to specify which data to select from a table. You use SQL to define complex queries. For example, you might select data from the `Employee` table to create a result that shows where each department is located, presenting the data sorted in increasing order by department number. This result is shown in Fig. 28.2. SQL queries are discussed in Section 28.4.

Department	Location
413	New Jersey
611	Orlando
642	Los Angeles

Fig. 28.2 | Result of selecting distinct `Department` and `Location` data from table `Employee`.

28.3 Relational Database Overview: The books Database

We now overview relational databases in the context of a sample `books` database we created for this chapter. Before we discuss SQL, we discuss the tables of the `books` database. We use this database to introduce various database concepts, including how to use SQL to obtain information from the database and to manipulate the data. We provide a script to create the database. You can find the script in the examples directory for this chapter. Section 28.7 explains how to use this script.

The database consists of three tables: `Authors`, `AuthorISBN` and `Titles`. The `Authors` table (described in Fig. 28.3) consists of three columns that maintain each author's unique ID number, first name and last name. Figure 28.4 contains sample data from the `Authors` table of the `books` database.

Column	Description
<code>AuthorID</code>	Author's ID number in the database. In the <code>books</code> database, this integer column is defined as autoincremented —for each row inserted in this table, the <code>AuthorID</code> value is increased by 1 automatically to ensure that each row has a unique <code>AuthorID</code> . This column represents the table's primary key.
<code>FirstName</code>	Author's first name (a string).
<code>LastName</code>	Author's last name (a string).

Fig. 28.3 | `Authors` table from the `books` database.

AuthorID	FirstName	LastName
1	Harvey	Deitel
2	Paul	Deitel
3	Andrew	Goldberg
4	David	Choffnes

Fig. 28.4 | Sample data from the Authors table.

The AuthorISBN table (described in Fig. 28.5) consists of two columns that maintain each ISBN and the corresponding author's ID number. This table associates authors with their books. Both columns are foreign keys that represent the relationship between the tables Authors and Titles—one row in table Authors may be associated with many rows in table Titles, and vice versa. The combined columns of the AuthorISBN table represent the table's primary key—thus, each row in this table must be a unique combination of an AuthorID and an ISBN. Figure 28.6 contains sample data from the AuthorISBN table of the books database. [Note: To save space, we have split the contents of this table into two columns, each containing the AuthorID and ISBN columns.] The AuthorID column is a **foreign key**—a column in this table that matches the primary key column in another table (i.e., AuthorID in the Authors table). Foreign keys are specified when creating a table. The foreign key helps maintain the **Rule of Referential Integrity**—every foreign-key value must appear as another table's primary-key value. This enables the DBMS to determine whether the AuthorID value for a particular book is valid. Foreign keys also allow related data in multiple tables to be selected from those tables for analytic purposes—this is known as **joining** the data.

Column	Description
AuthorID	The author's ID number, a foreign key to the Authors table.
ISBN	The ISBN for a book, a foreign key to the Titles table.

Fig. 28.5 | AuthorISBN table from the books database.

AuthorID	ISBN	AuthorID	ISBN
1	0131869000	2	0131450913
2	0131869000	1	0131828274
1	0132222205	2	0131828274
2	0132222205	3	0131450913
1	0131450913	4	0131828274

Fig. 28.6 | Sample data from the AuthorISBN table of books.

The Titles table described in Fig. 28.7 consists of four columns that stand for the ISBN, the title, the edition number and the copyright year. The table is in Fig. 28.8.

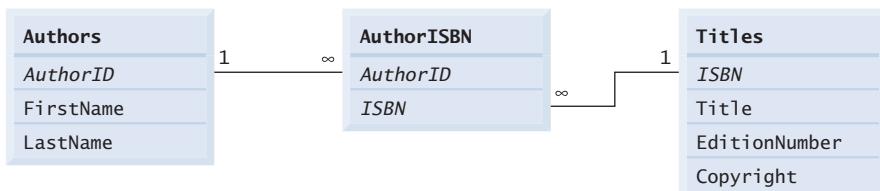
Column	Description
ISBN	ISBN of the book (a string). The table's primary key. ISBN is an abbreviation for "International Standard Book Number"—a numbering scheme that publishers use to give every book a unique identification number.
Title	Title of the book (a string).
EditionNumber	Edition number of the book (an integer).
Copyright	Copyright year of the book (a string).

Fig. 28.7 | Titles table from the books database.

ISBN	Title	EditionNumber	Copyright
0131869000	Visual Basic 2005 How to Program	3	2006
0131525239	Visual C# 2005 How to Program	2	2006
0132222205	Java How to Program	7	2007
0131857576	C++ How to Program	5	2005
0132404168	C How to Program	5	2007
0131450913	Internet & World Wide Web How to Program	3	2004
0131828274	Operating Systems	3	2004

Fig. 28.8 | Sample data from the Titles table of the books database.

There is a one-to-many relationship between a primary key and a corresponding foreign key (e.g., one publisher can publish many books). A foreign key can appear many times in its own table, but can appear only once (as the primary key) in another table. Figure 28.9 is an **entity-relationship (ER) diagram** for the books database. This diagram shows the database tables and the relationships among them. The first compartment in each box contains the table's name and the remaining compartments contain the table's columns. The names in italic are primary keys. A table's primary key uniquely identifies each row in the table. Every row must have a primary-key value, and that value must be unique in the table. This is known as the **Rule of Entity Integrity**. Again, for the AuthorISBN table, the primary key is the combination of both columns.

**Fig. 28.9** | Table relationships in the books database.

**Common Programming Error 28.1**

Not providing a value for every column in a primary key breaks the Rule of Entity Integrity and causes the DBMS to report an error.

**Common Programming Error 28.2**

Providing the same primary-key value in multiple rows causes the DBMS to report an error.

The lines connecting the tables (Fig. 28.9) represent the relationships between the tables. Consider the line between the AuthorISBN and Authors tables. On the Authors end of the line is a 1, and on the AuthorISBN end is an infinity symbol (∞), indicating a **one-to-many relationship** in which every author in the Authors table can have an arbitrary number of books in the AuthorISBN table. The relationship line links the AuthorID column in Authors (i.e., its primary key) to the AuthorID column in AuthorISBN (i.e., its foreign key). The AuthorID column in the AuthorISBN table is a foreign key.

**Common Programming Error 28.3**

Providing a foreign-key value that does not appear as a primary-key value in another table breaks the Rule of Referential Integrity and causes the DBMS to report an error.

The line between Titles and AuthorISBN illustrates another one-to-many relationship; a title can be written by any number of authors. In fact, the sole purpose of the AuthorISBN table is to provide a many-to-many relationship between Authors and Titles—an author can write many books and a book can have many authors.

28.4 SQL

We now provide an overview of SQL in the context of our books database. You'll be able to use the SQL discussed here in the examples later in the chapter and in examples in Chapters 30–31.

The next several subsections discuss the SQL keywords listed in Fig. 28.10 in the context of SQL queries and statements. Other SQL keywords are beyond this text's scope. To learn other keywords, refer to the SQL reference guide supplied by the vendor of the RDBMS you are using. [Note: For more information on SQL, refer to the web resources in Section 28.15.]

SQL keyword	Description
SELECT	Retrieves data from one or more tables.
FROM	Tables involved in the query. Required in every SELECT.
WHERE	Criteria for selection that determine the rows to be retrieved, deleted or updated. Optional in a SQL query or a SQL statement.
GROUP BY	Criteria for grouping rows. Optional in a SELECT query.
ORDER BY	Criteria for ordering rows. Optional in a SELECT query.

Fig. 28.10 | SQL query keywords. (Part 1 of 2.)

SQL keyword	Description
INNER JOIN	Merge rows from multiple tables.
INSERT	Insert rows into a specified table.
UPDATE	Update rows in a specified table.
DELETE	Delete rows from a specified table.

Fig. 28.10 | SQL query keywords. (Part 2 of 2.)

28.4.1 Basic SELECT Query

Let us consider several SQL queries that extract information from database books. A SQL query “selects” rows and columns from one or more tables in a database. Such selections are performed by queries with the **SELECT** keyword. The basic form of a **SELECT** query is

```
SELECT * FROM tableName
```

in which the **asterisk (*)** wildcard character indicates that all columns from the *tableName* table should be retrieved. For example, to retrieve all the data in the *Authors* table, use

```
SELECT * FROM Authors
```

Most programs do not require all the data in a table. To retrieve only specific columns, replace the asterisk (*) with a comma-separated list of the column names. For example, to retrieve only the columns *AuthorID* and *LastName* for all rows in the *Authors* table, use the query

```
SELECT AuthorID, LastName FROM Authors
```

This query returns the data listed in Fig. 28.11.

AuthorID	LastName
1	Deitel
2	Deitel
3	Goldberg
4	Choffnes

Fig. 28.11 | Sample *AuthorID* and *LastName* data from the *Authors* table.

Software Engineering Observation 28.3

For most queries, the asterisk (*) should not be used to specify column names. In general, you process results by knowing in advance the order of the columns in the result—for example, selecting *AuthorID* and *LastName* from table *Authors* ensures that the columns will appear in the result with *AuthorID* as the first column and *LastName* as the second column. Programs typically process result columns by specifying the column number in the result (starting from number 1 for the first column). Selecting columns by name also avoids returning unneeded columns and protects against changes in the actual order of the columns in the table(s) by returning the columns in the exact order specified.



Common Programming Error 28.4

If you assume that the columns are always returned in the same order from a query that uses the asterisk (*), the program may process the results incorrectly. If the column order in the table(s) changes or if additional columns are added at a later time, the order of the columns in the result will change accordingly.

28.4.2 WHERE Clause

In most cases, it is necessary to locate rows in a database that satisfy certain **selection criteria**. Only rows that satisfy the selection criteria (formally called **predicates**) are selected. SQL uses the optional **WHERE clause** in a query to specify the selection criteria for the query. The basic form of a query with selection criteria is

```
SELECT columnName1, columnName2, ... FROM tableName WHERE criteria
```

For example, to select the **Title**, **EditionNumber** and **Copyright** columns from table **Titles** for which the **Copyright** date is greater than 2005, use the query

```
SELECT Title, EditionNumber, Copyright  
FROM Titles  
WHERE Copyright > '2005'
```

Notice that strings in SQL are delimited by single ('') rather than double ("") quotes.

Figure 28.12 shows the result of the preceding query. The **WHERE** clause criteria can contain the operators <, >, <=, >=, =, <> and **LIKE**. Operator **LIKE** is used for **pattern matching** with wildcard characters **percent (%)** and **underscore (_)**. Pattern matching allows SQL to search for strings that match a given pattern.

Title	EditionNumber	Copyright
Visual C# 2005 How to Program	2	2006
Visual Basic 2005 How to Program	3	2006
Java How to Program	7	2007
C How to Program	5	2007

Fig. 28.12 | Sampling of titles with copyrights after 2005 from table **Titles**.

A pattern that contains a percent character (%) searches for strings that have zero or more characters at the percent character's position in the pattern. For example, the next query locates the rows of all the authors whose last name starts with the letter D:

```
SELECT AuthorID, FirstName, LastName  
FROM Authors  
WHERE LastName LIKE 'D%'
```

This query selects the two rows shown in Fig. 28.13—two of the four authors have a last name starting with the letter D (followed by zero or more characters). The % symbol in the **WHERE** clause's **LIKE** pattern indicates that any number of characters can appear after the letter D in the **LastName**. Note that the pattern string is surrounded by single-quote characters.

AuthorID	FirstName	LastName
1	Harvey	Deitel
2	Paul	Deitel

Fig. 28.13 | Authors whose last name starts with D from the Authors table.



Portability Tip 28.1

See the documentation for your database system to determine whether SQL is case sensitive on your system and to determine the syntax for SQL keywords (i.e., should they be all uppercase letters, all lowercase letters or some combination of the two?).



Portability Tip 28.2

Read your database system's documentation carefully to determine whether it supports the LIKE operator as discussed here. The SQL we discuss is supported by most RDBMSs, but it is always a good idea to check the features of SQL that are supported by your RDBMS.

An underscore (_) in the pattern string indicates a single wildcard character at that position in the pattern. For example, the following query locates the rows of all the authors whose last names start with any character (specified by _), followed by the letter o, followed by any number of additional characters (specified by %):

```
SELECT AuthorID, FirstName, LastName
  FROM Authors
 WHERE LastName LIKE '_o%'
```

The preceding query produces the row shown in Fig. 28.14, because only one author in our database has a last name that contains the letter o as its second letter.

AuthorID	FirstName	LastName
3	Andrew	Goldberg

Fig. 28.14 | The only author from the Authors table whose last name contains o as the second letter.

28.4.3 ORDER BY Clause

The rows in the result of a query can be sorted into ascending or descending order by using the optional **ORDER BY clause**. The basic form of a query with an ORDER BY clause is

```
SELECT columnName1, columnName2, ... FROM tableName ORDER BY column ASC
SELECT columnName1, columnName2, ... FROM tableName ORDER BY column DESC
```

where ASC specifies ascending order (lowest to highest), DESC specifies descending order (highest to lowest) and *column* specifies the column on which the sort is based. For example, to obtain the list of authors in ascending order by last name (Fig. 28.15), use the query

```
SELECT AuthorID, FirstName, LastName
  FROM Authors
 ORDER BY LastName ASC
```

AuthorID	FirstName	LastName
4	David	Choffnes
1	Harvey	Deitel
2	Paul	Deitel
3	Andrew	Goldberg

Fig. 28.15 | Sample data from table Authors in ascending order by LastName.

Note that the default sorting order is ascending, so ASC is optional. To obtain the same list of authors in descending order by last name (Fig. 28.16), use the query

```
SELECT AuthorID, FirstName, LastName
  FROM Authors
 ORDER BY LastName DESC
```

AuthorID	FirstName	LastName
3	Andrew	Goldberg
1	Harvey	Deitel
2	Paul	Deitel
4	David	Choffnes

Fig. 28.16 | Sample data from table Authors in descending order by LastName.

Multiple columns can be used for sorting with an ORDER BY clause of the form

```
ORDER BY column1 sortOrder, column2 sortOrder, ...
```

where *sortingOrder* is either ASC or DESC. Note that the *sortingOrder* does not have to be identical for each column. The query

```
SELECT AuthorID, FirstName, LastName
  FROM Authors
 ORDER BY LastName, FirstName
```

sorts all the rows in ascending order by last name, then by first name. If any rows have the same last-name value, they're returned sorted by first name (Fig. 28.17).

AuthorID	FirstName	LastName
4	David	Choffnes
1	Harvey	Deitel
2	Paul	Deitel
3	Andrew	Goldberg

Fig. 28.17 | Sample data from Authors in ascending order by LastName and FirstName.

The WHERE and ORDER BY clauses can be combined in one query, as in

```
SELECT ISBN, Title, EditionNumber, Copyright
  FROM Titles
 WHERE Title LIKE '%How to Program'
 ORDER BY Title ASC
```

which returns the ISBN, Title, EditionNumber and Copyright of each book in the Titles table that has a Title ending with "How to Program" and sorts them in ascending order by Title. The query results are shown in Fig. 28.18.

ISBN	Title	Edition-Number	Copy-right
0132404168	C How to Program	5	2007
0131857576	C++ How to Program	5	2005
0131450913	Internet and World Wide Web How to Program	3	2004
0132222205	Java How to Program	7	2007
0131869000	Visual Basic 2005 How to Program	3	2006
013152539	Visual C# 2005 How to Program	2	2006

Fig. 28.18 | Sampling of books from table Titles whose titles end with How to Program in ascending order by Title.

28.4.4 Merging Data from Multiple Tables: INNER JOIN

Database designers often split related data into separate tables to ensure that a database does not store data redundantly. For example, the books database has tables Authors and Titles. We use an AuthorISBN table to store the relationship data between authors and their corresponding titles. If we did not separate this information into individual tables, we would need to include author information with each entry in the Titles table. This would result in the database's storing duplicate author information for authors who wrote multiple books. Often, it is necessary to merge data from multiple tables into a single result. Referred to as joining the tables, this is specified by an **INNER JOIN** operator in the query. An INNER JOIN merges rows from two tables by matching values in columns that are common to the tables. The basic form of an INNER JOIN is:

```
SELECT columnName1, columnName2, ...
  FROM table1
 INNER JOIN table2
    ON table1.columnName = table2.columnName
```

The **ON clause** of the INNER JOIN specifies the columns from each table that are compared to determine which rows are merged. For example, the following query produces a list of authors accompanied by the ISBNs for books written by each author:

```
SELECT FirstName, LastName, ISBN
  FROM Authors
 INNER JOIN AuthorISBN
    ON Authors.AuthorID = AuthorISBN.AuthorID
 ORDER BY LastName, FirstName
```

The query merges the `FirstName` and `LastName` columns from table `Authors` with the `ISBN` column from table `AuthorISBN`, sorting the result in ascending order by `LastName` and `FirstName`. Note the use of the syntax `tableName.columnName` in the `ON` clause. This syntax, called a **qualified name**, specifies the columns from each table that should be compared to join the tables. The “`tableName.`” syntax is required if the columns have the same name in both tables. The same syntax can be used in any SQL statement to distinguish columns in different tables that have the same name. In some systems, table names qualified with the database name can be used to perform cross-database queries. As always, the query can contain an `ORDER BY` clause. Figure 28.19 shows the results of the preceding query, ordered by `LastName` and `FirstName`. [Note: To save space, we split the result of the query into two columns, each containing the `FirstName`, `LastName` and `ISBN` columns.]

FirstName	LastName	ISBN	FirstName	LastName	ISBN
David	Choffnes	0131828274	Paul	Deitel	0131869000
Harvey	Deitel	0131869000	Paul	Deitel	0131525239
Harvey	Deitel	0131525239	Paul	Deitel	0132222205
Harvey	Deitel	0132222205	Paul	Deitel	0131857576
Harvey	Deitel	0131857576	Paul	Deitel	0132404168
Harvey	Deitel	0132404168	Paul	Deitel	0131450913
Harvey	Deitel	0131450913	Paul	Deitel	0131828274
Harvey	Deitel	0131828274	Andrew	Goldberg	0131450913

Fig. 28.19 | Sampling of authors and ISBNs for the books they have written in ascending order by `LastName` and `FirstName`.



Software Engineering Observation 28.4

If a SQL statement includes columns with the same name from multiple tables, the statement must precede those column names with their table names and a dot (e.g., `Authors.AuthorID`).



Common Programming Error 28.5

Failure to qualify names for columns that have the same name in two or more tables is an error.

28.4.5 INSERT Statement

The `INSERT` statement inserts a row into a table. The basic form of this statement is

```
INSERT INTO tableName ( columnName1, columnName2, ..., columnNameN )
VALUES ( value1, value2, ..., valueN )
```

where `tableName` is the table in which to insert the row. The `tableName` is followed by a comma-separated list of column names in parentheses (this list is not required if the `INSERT` operation specifies a value for every column of the table in the correct order). The list of column names is followed by the SQL keyword `VALUES` and a comma-separated list of values in parentheses. The values specified here must match the columns specified after the table name in both order and type (e.g., if `columnName1` is supposed to be the `FirstName`

column, then *value1* should be a string in single quotes representing the first name). Always explicitly list the columns when inserting rows. If the table's column order changes or a new column is added, using only VALUES may cause an error. The INSERT statement

```
INSERT INTO Authors ( FirstName, LastName )
VALUES ( 'Sue', 'Smith' )
```

inserts a row into the Authors table. The statement indicates that values are provided for the FirstName and LastName columns. The corresponding values are 'Sue' and 'Smith'. We do not specify an AuthorID in this example because AuthorID is an autoincremented column in the Authors table. For every row added to this table, the DBMS assigns a unique AuthorID value that is the next value in the autoincremented sequence (i.e., 1, 2, 3 and so on). In this case, Sue Smith would be assigned AuthorID number 5. Figure 28.20 shows the Authors table after the INSERT operation. [Note: Not every database management system supports autoincremented columns. Check the documentation for your DBMS for alternatives to autoincremented columns.]

AuthorID	FirstName	LastName
1	Harvey	Deitel
2	Paul	Deitel
3	Andrew	Goldberg
4	David	Choffnes
5	Sue	Smith

Fig. 28.20 | Sample data from table Authors after an INSERT operation.



Common Programming Error 28.6

It is normally an error to specify a value for an autoincrement column.



Common Programming Error 28.7

SQL delimits strings with single quotes (''). A string containing a single quote (e.g., O'Malley) must have two single quotes in the position where the single quote appears (e.g., 'O''MalleY'). The first acts as an escape character for the second. Not escaping single-quote characters in a string that is part of a SQL statement is a SQL syntax error.

28.4.6 UPDATE Statement

An UPDATE statement modifies data in a table. Its basic form is

```
UPDATE tableName
SET columnName1 = value1, columnName2 = value2, ..., columnNameN = valueN
WHERE criteria
```

where *tableName* is the table to update. The *tableName* is followed by keyword **SET** and a comma-separated list of column name/value pairs in the format *columnName* = *value*. The optional **WHERE** clause provides criteria that determine which rows to update. Though not required, the **WHERE** clause is typically used, unless a change is to be made to every row. The UPDATE statement

```
UPDATE Authors
  SET LastName = 'Jones'
WHERE LastName = 'Smith' AND FirstName = 'Sue'
```

updates a row in the Authors table. The statement indicates that LastName will be assigned the value Jones for the row in which LastName is equal to Smith and FirstName is equal to Sue. [Note: If there are multiple rows with the first name “Sue” and the last name “Smith,” this statement will modify all such rows to have the last name “Jones.”] If we know the AuthorID in advance of the UPDATE operation (possibly because we searched for it previously), the WHERE clause can be simplified as follows:

```
WHERE AuthorID = 5
```

Figure 28.21 shows the Authors table after the UPDATE operation has taken place.

AuthorID	FirstName	LastName
1	Harvey	Deitel
2	Paul	Deitel
3	Andrew	Goldberg
4	David	Choffnes
5	Sue	Jones

Fig. 28.21 | Sample data from table Authors after an UPDATE operation.

28.4.7 DELETE Statement

A SQL **DELETE** statement removes rows from a table. Its basic form is

```
DELETE FROM tableName WHERE criteria
```

where *tableName* is the table from which to delete. The optional WHERE clause specifies the criteria used to determine which rows to delete. If this clause is omitted, all the table’s rows are deleted. The DELETE statement

```
DELETE FROM Authors
  WHERE LastName = 'Jones' AND FirstName = 'Sue'
```

deletes the row for Sue Jones in the Authors table. If we know the AuthorID in advance of the DELETE operation, the WHERE clause can be simplified as follows:

```
WHERE AuthorID = 5
```

Figure 28.22 shows the Authors table after the DELETE operation has taken place.

AuthorID	FirstName	LastName
1	Harvey	Deitel
2	Paul	Deitel
3	Andrew	Goldberg
4	David	Choffnes

Fig. 28.22 | Sample data from table Authors after a DELETE operation.

28.5 Instructions for Installing MySQL and MySQL Connector/J

MySQL 5.0 Community Edition is an open-source database management system that executes on many platforms, including Windows, Solaris, Linux, and Macintosh. Complete information about MySQL is available from www.mysql.com. The examples in Sections 28.8—28.9 manipulate MySQL databases.

Installing MySQL

To install MySQL Community Edition:

1. To learn about the installation requirements for your platform, visit the site dev.mysql.com/doc/refman/5.0/en/installing-cs.html.
2. Visit dev.mysql.com/downloads/mysql/5.0.html and download the installer for your platform. For the MySQL examples in this chapter, you need only the Windows Essentials package on Microsoft Windows, or the Standard package on most other platforms. [Note: For these instructions, we assume you are running Microsoft Windows. Complete installation instructions for other platforms are available at dev.mysql.com/doc/refman/5.0/en/installing.html.]
3. Double click `mysql-essential-5.0.67-win32.msi` to start the installer. [Note: This name may differ based on the current version of MySQL 5.0.] Click **Next >**.
4. Choose **Typical** for the **Setup Type** and click **Next >**. Then click **Install**.

When the installation completes, click **Next >** twice, then **Finish** to begin configuring the server. To configure the server:

1. Click **Next >**, then select **Standard Configuration** and click **Next >** again.
2. You have the option of installing MySQL as a Windows service, which enables the MySQL server to begin executing automatically each time your system starts. For our examples, this is unnecessary, so uncheck **Install as a Windows Service**, then check **Include Bin Directory in Windows PATH**. This will enable you to use the MySQL commands in the Windows Command Prompt.
3. Click **Next >**, then click **Execute** to perform the server configuration.
4. Click **Finish** to close the wizard.

Installing MySQL Connector/J

To use MySQL with JDBC, you also need to install **MySQL Connector/J** (the J stands for Java)—a JDBC driver that allows programs to use JDBC to interact with MySQL. MySQL Connector/J can be downloaded from

dev.mysql.com/downloads/connector/j/5.1.html

The documentation for Connector/J is located at dev.mysql.com/doc/connector/j/en/connector-j.html. At the time of this writing, the current generally available release of MySQL Connector/J is 5.1.7. To install MySQL Connector/J:

1. Download `mysql-connector-java-5.1.7.tar.gz`.
2. Open `mysql-connector-java-5.1.7.tar.gz` with a file extractor, such as WinZip (www.winzip.com). Extract its contents to the C:\ drive. This will create a directory named `mysql-connector-java-5.1.7`. This folder's `docs` subdirectory

contains the MySQL Connector/J documentation (`connector-j.pdf`). You can also view it online at dev.mysql.com/doc/connector/j/en/connector-j.html.

28.6 Instructions for Setting Up a MySQL User Account

For the MySQL examples to execute correctly, you need to set up a user account that allows users to create, delete and modify a database. After MySQL is installed, follow the steps below to set up a user account (these steps assume MySQL is installed in its default installation directory):

1. Open a Command Prompt and start the database server by executing the command `mysqld-nt.exe`. Note that this command has no output—it simply starts the MySQL server. Do not close this window—doing so terminates the server.
2. Next, you'll start the MySQL monitor so you can set up a user account, open another Command Prompt and execute the command

```
mysql -h localhost -u root
```

The `-h` option indicates the host (i.e., computer) on which the MySQL server is running—in this case your local computer (`localhost`). The `-u` option indicates the user account that will be used to log in to the server—`root` is the default user account that is created during installation to allow you to configure the server. Once you've logged in, you'll see a `mysql>` prompt at which you can type commands to interact with the MySQL server.

3. At the `mysql>` prompt, type

```
USE mysql;
```

to select the built-in database named `mysql`, which stores server information, such as user accounts and their privileges for interacting with the server. Note that each command must end with a semicolon. To confirm the command, MySQL issues the message “Database changed.”

4. Next, you'll add the `deitel` user account to the `mysql` built-in database. The `mysql` database contains a table called `user` with columns that represent the user's name, password and various privileges. To create the `deitel` user account with the password `deitel`, execute the following commands from the `mysql>` prompt:

```
create user 'deitel'@'localhost' identified by 'deitel';
grant select, insert, update, delete, create, drop, references,
execute on *.* to 'deitel'@'localhost';
```

This creates the `deitel` user with the privileges needed to create the databases used in this chapter and manipulate them.

5. Type the command

```
exit;
```

to terminate the MySQL monitor.

28.7 Creating Database books in MySQL

For each MySQL database we discuss, we provide a SQL script in a `.sql` file that sets up the database and its tables. You can execute these scripts in the MySQL monitor. In this

chapter's examples directory, you'll find the script `books.sql` to create the `books` database. For the following steps, we assume that the MySQL server (`mysqld-nt.exe`) is still running. To execute the `books.sql` script:

1. Open a Command Prompt and use the `cd` command to change directories to the location that contains the `books.sql` script.
2. Start the MySQL monitor by typing

```
mysql -h localhost -u deitel -p
```

The `-p` option prompts you for the password for the `deitel` user account. When prompted, enter the password `deitel`.

3. Execute the script by typing

```
source books.sql;
```

This creates a new directory named `books` in the server's data directory—located by default on Windows at `C:\Program Files\MySQL\MySQL Server 5.0\data`. This new directory contains the `books` database.

4. Type the command

```
exit;
```

to terminate the MySQL monitor. You are now ready to proceed to the first JDBC example.

28.8 Manipulating Databases with JDBC

In this section, we present two examples. The first introduces how to connect to a database and query the database. The second demonstrates how to display the result of the query in a `JTable`.

28.8.1 Connecting to and Querying a Database

The example of Fig. 28.23 performs a simple query on the `books` database that retrieves the entire `Authors` table and displays the data. The program illustrates connecting to the database, querying the database and processing the result. The discussion that follows presents the key JDBC aspects of the program. [Note: Sections 28.5—28.7 demonstrate how to start the MySQL server, configure a user account and create the `books` database. These steps *must* be performed before executing the program of Fig. 28.23.]

```

1 // Fig. 28.23: DisplayAuthors.java
2 // Displaying the contents of the Authors table.
3 import java.sql.Connection;
4 import java.sql.Statement;
5 import java.sql.DriverManager;
6 import java.sql.ResultSet;
7 import java.sql.ResultSetMetaData;
8 import java.sql.SQLException;
9

```

Fig. 28.23 | Displaying the contents of the `Authors` table. (Part I of 3.)

```
10 public class DisplayAuthors
11 {
12     // database URL
13     static final String DATABASE_URL = "jdbc:mysql://localhost/books";
14
15     // launch the application
16     public static void main( String args[] )
17     {
18         Connection connection = null; // manages connection
19         Statement statement = null; // query statement
20         ResultSet resultSet = null; // manages results
21
22         // connect to database books and query database
23         try
24         {
25             // establish connection to database
26             connection = DriverManager.getConnection(
27                 DATABASE_URL, "deitel", "deitel" );
28
29             // create Statement for querying database
30             statement = connection.createStatement();
31
32             // query database
33             resultSet = statement.executeQuery(
34                 "SELECT AuthorID, FirstName, LastName FROM Authors" );
35
36             // process query results
37             ResultSetMetaData metaData = resultSet.getMetaData();
38             int numberOfColumns = metaData.getColumnCount();
39             System.out.println( "Authors Table of Books Database:\n" );
40
41             for ( int i = 1; i <= numberOfColumns; i++ )
42                 System.out.printf( "%-8s\t", metaData.getColumnName( i ) );
43             System.out.println();
44
45             while ( resultSet.next() )
46             {
47                 for ( int i = 1; i <= numberOfColumns; i++ )
48                     System.out.printf( "%-8s\t", resultSet.getObject( i ) );
49                 System.out.println();
50             } // end while
51         } // end try
52         catch ( SQLException sqlException )
53         {
54             sqlException.printStackTrace();
55         } // end catch
56         finally // ensure resultSet, statement and connection are closed
57         {
58             try
59             {
60                 resultSet.close();
61                 statement.close();
62                 connection.close();
63             } // end try

```

Fig. 28.23 | Displaying the contents of the Authors table. (Part 2 of 3.)

```

64     catch ( Exception exception )
65     {
66         exception.printStackTrace();
67     } // end catch
68 } // end finally
69 } // end main
70 } // end class DisplayAuthors

```

Authors Table of Books Database:

AuthorID	FirstName	LastName
1	Harvey	Deitel
2	Paul	Deitel
3	Andrew	Goldberg
4	David	Choffnes

Fig. 28.23 | Displaying the contents of the Authors table. (Part 3 of 3.)

Lines 3–8 import the JDBC interfaces and classes from package `java.sql` used in this program. Line 13 declares a string constant for the database URL. This identifies the name of the database to connect to, as well as information about the protocol used by the JDBC driver (discussed shortly). Method `main` (lines 16–69) connects to the books database, queries the database, displays the result of the query and closes the database connection.

In past versions of Java, programs were required to load an appropriate database driver before connecting to a database. JDBC 4.0, part of Java SE 6, supports **automatic driver discovery**—you are no longer required to load the database driver in advance. To ensure that the program can locate the database driver class, you must include the class's location in the program's classpath when you execute the program. For MySQL, you include the file `mysql-connector-java-5.1.7-bin.jar` (in the `C:\mysql-connector-java-5.1.7` directory) in your program's classpath, as in:

```

java -classpath
    .;c:\mysql-connector-java-5.1.7\mysql-connector-java-5.1.7-bin.jar
DisplayAuthors

```

Notice the period (.) at the beginning of the classpath information. If it is missing, the JVM will not look for classes in the current directory and thus will not find the `DisplayAuthors` class file. You may also copy the `mysql-connector-java-5.1.7-bin.jar` file to your JDK's `\jre\lib\ext` folder. After doing so, you can run the application simply using the command

```
java DisplayAuthors
```

Lines 26–27 of Fig. 28.23 create a `Connection` object (package `java.sql`) referenced by `connection`. An object that implements interface `Connection` manages the connection between the Java program and the database. `Connection` objects enable programs to create SQL statements that manipulate databases. The program initializes `connection` with the result of a call to static method `getConnection` of class `DriverManager` (package `java.sql`), which attempts to connect to the database specified by its URL. Method `get-Connection` takes three arguments—a `String` that specifies the database URL, a `String` that specifies the username and a `String` that specifies the password. The username and password are set in Section 28.6. If you used a different username and password, you need

to replace the username (second argument) and password (third argument) passed to method `getConnection` in line 27. The URL locates the database (possibly on a network or in the local file system of the computer). The URL `jdbc:mysql://localhost/books` specifies the protocol for communication (`jdbc`), the **subprotocol** for communication (`mysql`) and the location of the database (`//localhost/books`, where `localhost` is the host running the MySQL server and `books` is the database name). The subprotocol `mysql` indicates that the program uses a MySQL-specific subprotocol to connect to the MySQL database. If the `DriverManager` cannot connect to the database, method `getConnection` throws a **`SQLException`** (package `java.sql`). Figure 28.24 lists the JDBC driver names and database URL formats of several popular RDBMSs.

RDBMS	Database URL format
MySQL	<code>jdbc:mysql://hostname:portNumber/databaseName</code>
ORACLE	<code>jdbc:oracle:thin:@hostname:portNumber:databaseName</code>
DB2	<code>jdbc:db2:hostname:portNumber/databaseName</code>
Java DB/Apache	<code>jdbc:derby:DataBaseName</code> (embedded)
Derby	<code>jdbc:derby://hostname:portNumber/databaseName</code> (network)
Microsoft SQL Server	<code>jdbc:sqlserver://hostname:portNumber;databaseName=DataBaseName</code>
Sybase	<code>jdbc:sybase:Tds:hostname:portNumber/databaseName</code>

Fig. 28.24 | Popular JDBC database URL formats.



Software Engineering Observation 28.5

Most database management systems require the user to log in before accessing the database contents. `DriverManager` method `getConnection` is overloaded with versions that enable the program to supply the user name and password to gain access.

Line 30 invokes `Connection` method `createStatement` to obtain an object that implements interface `Statement` (package `java.sql`). The program uses the `Statement` object to submit SQL statements to the database.

Lines 33–34 use the `Statement` object's `executeQuery` method to submit a query that selects all the author information from table `Authors`. This method returns an object that implements interface `ResultSet` and contains the query results. The `ResultSet` methods enable the program to manipulate the query result.

Lines 37–50 process the `ResultSet`. Line 37 obtains the metadata for the `ResultSet` as a `ResultSetMetaData` (package `java.sql`) object. The `metadata` describes the `ResultSet`'s contents. Programs can use metadata programmatically to obtain information about the `ResultSet`'s column names and types. Line 38 uses `ResultSetMetaData` method `getColumnCount` to retrieve the number of columns in the `ResultSet`. Lines 41–42 display the column names.



Software Engineering Observation 28.6

Metadata enables programs to process `ResultSet` contents dynamically when detailed information about the `ResultSet` is not known in advance.

Lines 45–50 display the data in each `ResultSet` row. First, the program positions the `ResultSet` cursor (which points to the row being processed) to the first row in the `ResultSet` with method `next` (line 45). Method `next` returns boolean value `true` if it is able to position to the next row; otherwise, the method returns `false`.



Common Programming Error 28.8

Initially, a `ResultSet` cursor is positioned before the first row. A `SQLException` occurs if you attempt to access a `ResultSet`'s contents before positioning the `ResultSet` cursor to the first row with method `next`.

If there are rows in the `ResultSet`, lines 47–48 extract and display the contents of each column in the current row. When a `ResultSet` is processed, each column can be extracted as a specific Java type. In fact, `ResultSetMetaData` method `getColumnType` returns a constant integer from class `Types` (package `java.sql`) indicating the type of a specified column. Programs can use these values in a `switch` statement to invoke `ResultSet` methods that return the column values as appropriate Java types. If the type of a column is `Types.INTEGER`, `ResultSet` method `getInt` returns the column value as an `int`. `ResultSet` `get` methods typically receive as an argument either a column number (as an `int`) or a column name (as a `String`) indicating which column's value to obtain. Visit

```
java.sun.com/javase/6/docs/technotes/guides/jdbc/getstart/  
GettingStartedTOC.fm.html
```

for detailed mappings of SQL data types to Java types and to determine the appropriate `ResultSet` method to call for each SQL data type.



Performance Tip 28.1

If a query specifies the exact columns to select from the database, the `ResultSet` contains the columns in the specified order. In this case, using the column number to obtain the column's value is more efficient than using the column name. The column number provides direct access to the specified column. Using the column name requires a search of the column names to locate the appropriate column.



Error-Prevention Tip 28.1

Using column names to obtain values from a `ResultSet` produces code that is less error prone than obtaining values by column number—you don't need to remember the column order. Also, if the column order changes, your code does not have to change.

For simplicity, this example treats each value as an `Object`. The program retrieves each column value with `ResultSet` method `getObject` (line 48) and prints the `String` representation of the `Object`. Note that, unlike array indices, which start at 0, `ResultSet` column numbers start at 1. The `finally` block (lines 56–68) closes the `ResultSet` (line 60), the `Statement` (line 61) and the database `Connection` (line 62). [Note: Lines 60–62 will throw `NullPointerExceptions` if the `ResultSet`, `Statement` or `Connection` objects were not created properly. For code used in industry, you should check the variables that refer to these objects to see if they're `null` before you call `close`.]



Common Programming Error 28.9

Specifying column 0 when obtaining values from a `ResultSet` causes a `SQLException`.



Common Programming Error 28.10

A `SQLException` occurs if you attempt to manipulate a `ResultSet` after closing the `Statement` that created it. The `ResultSet` is discarded when the corresponding `Statement` is closed.



Software Engineering Observation 28.7

Each `Statement` object can open only one `ResultSet` object at a time. When a `Statement` returns a new `ResultSet`, the `Statement` closes the prior `ResultSet`. To use multiple `ResultSets` in parallel, separate `Statement` objects must return the `ResultSets`.

28.8.2 Querying the books Database

The next example (Fig. 28.25 and Fig. 28.28) allows the user to enter any query into the program. The example displays the result of a query in a `JTable`, using a `TableModel` object to provide the `ResultSet` data to the `JTable`. A `JTable` is a swing GUI component that can be bound to a database to display the results of a query. Class `ResultSetTableModel` (Fig. 28.25) performs the connection to the database via a `TableModel` and maintains the `ResultSet`. Class `DisplayQueryResults` (Fig. 28.28) creates the GUI and specifies an instance of class `ResultSetTableModel` to provide data for the `JTable`.

ResultSetTableModel Class

Class `ResultSetTableModel` (Fig. 28.25) extends class `AbstractTableModel` (package `javax.swing.table`), which implements interface `TableModel`. `ResultSetTableModel` overrides `TableModel` methods `getColumnClass`, `getRowCount`, `getColumnName`, `getRowCount` and `getValueAt`. The default implementations of `TableModel` methods `isCellEditable` and `setValueAt` (provided by `AbstractTableModel`) are not overridden, because this example does not support editing the `JTable` cells. The default implementations of `TableModel` methods `addTableModelListener` and `removeTableModelListener` (provided by `AbstractTableModel`) are not overridden, because the implementations of these methods in `AbstractTableModel` properly add and remove event listeners.

```
1 // Fig. 28.25: ResultSetTableModel.java
2 // A TableModel that supplies ResultSet data to a JTable.
3 import java.sql.Connection;
4 import java.sql.Statement;
5 import java.sql.DriverManager;
6 import java.sql.ResultSet;
7 import java.sql.ResultSetMetaData;
8 import java.sql.SQLException;
9 import javax.swing.table.AbstractTableModel;
10
11 // ResultSet rows and columns are counted from 1 and JTable
12 // rows and columns are counted from 0. When processing
13 // ResultSet rows or columns for use in a JTable, it is
14 // necessary to add 1 to the row or column number to manipulate
15 // the appropriate ResultSet column (i.e., JTable column 0 is
16 // ResultSet column 1 and JTable row 0 is ResultSet row 1).
17 public class ResultSetTableModel extends AbstractTableModel
18 {
```

Fig. 28.25 | A `TableModel` that supplies `ResultSet` data to a `JTable`. (Part 1 of 5.)

```
19  private Connection connection;
20  private Statement statement;
21  private ResultSet resultSet;
22  private ResultSetMetaData metaData;
23  private int numberOfRows;
24
25  // keep track of database connection status
26  private boolean connectedToDatabase = false;
27
28  // constructor initializes resultSet and obtains its meta data object;
29  // determines number of rows
30  public ResultSetTableModel( String url, String username,
31      String password, String query ) throws SQLException
32  {
33      // connect to database
34      connection = DriverManager.getConnection( url, username, password );
35
36      // create Statement to query database
37      statement = connection.createStatement(
38          ResultSet.TYPE_SCROLL_INSENSITIVE,
39          ResultSet.CONCUR_READ_ONLY );
40
41      // update database connection status
42      connectedToDatabase = true;
43
44      // set query and execute it
45      setQuery( query );
46  } // end constructor ResultSetTableModel
47
48  // get class that represents column type
49  public Class getColumnClass( int column ) throws IllegalStateException
50  {
51      // ensure database connection is available
52      if ( !connectedToDatabase )
53          throw new IllegalStateException( "Not Connected to Database" );
54
55      // determine Java class of column
56      try
57      {
58          String className = metaData.getColumnClassName( column + 1 );
59
60          // return Class object that represents className
61          return Class.forName( className );
62      } // end try
63      catch ( Exception exception )
64      {
65          exception.printStackTrace();
66      } // end catch
67
68      return Object.class; // if problems occur above, assume type Object
69  } // end method getColumnClass
70
```

Fig. 28.25 | A TableModel that supplies ResultSet data to a JTable. (Part 2 of 5.)

```
71 // get number of columns in ResultSet
72 public int getColumnCount() throws IllegalStateException
73 {
74     // ensure database connection is available
75     if ( !connectedToDatabase )
76         throw new IllegalStateException( "Not Connected to Database" );
77
78     // determine number of columns
79     try
80     {
81         return metaData.getColumnCount();
82     } // end try
83     catch ( SQLException sqlException )
84     {
85         sqlException.printStackTrace();
86     } // end catch
87
88     return 0; // if problems occur above, return 0 for number of columns
89 } // end method getColumnCount
90
91 // get name of a particular column in ResultSet
92 public String getColumnName( int column ) throws IllegalStateException
93 {
94     // ensure database connection is available
95     if ( !connectedToDatabase )
96         throw new IllegalStateException( "Not Connected to Database" );
97
98     // determine column name
99     try
100    {
101        return metaData.getColumnName( column + 1 );
102    } // end try
103    catch ( SQLException sqlException )
104    {
105        sqlException.printStackTrace();
106    } // end catch
107
108    return ""; // if problems, return empty string for column name
109 } // end method getColumnName
110
111 // return number of rows in ResultSet
112 public int getRowCount() throws IllegalStateException
113 {
114     // ensure database connection is available
115     if ( !connectedToDatabase )
116         throw new IllegalStateException( "Not Connected to Database" );
117
118     return numberOfRows;
119 } // end method getRowCount
120
121 // obtain value in particular row and column
122 public Object getValueAt( int row, int column )
123     throws IllegalStateException
124 {
```

Fig. 28.25 | A `TableModel` that supplies `ResultSet` data to a `JTable`. (Part 3 of 5.)

```
125     // ensure database connection is available
126     if ( !connectedToDatabase )
127         throw new IllegalStateException( "Not Connected to Database" );
128
129     // obtain a value at specified ResultSet row and column
130     try
131     {
132         resultSet.absolute( row + 1 );
133         return resultSet.getObject( column + 1 );
134     } // end try
135     catch ( SQLException sqlException )
136     {
137         sqlException.printStackTrace();
138     } // end catch
139
140     return ""; // if problems, return empty string object
141 } // end method getValueAt
142
143 // set new database query string
144 public void setQuery( String query )
145     throws SQLException, IllegalStateException
146 {
147     // ensure database connection is available
148     if ( !connectedToDatabase )
149         throw new IllegalStateException( "Not Connected to Database" );
150
151     // specify query and execute it
152     resultSet = statement.executeQuery( query );
153
154     // obtain meta data for ResultSet
155     metaData = resultSet.getMetaData();
156
157     // determine number of rows in ResultSet
158     resultSet.last(); // move to last row
159     numberOfRows = resultSet.getRow(); // get row number
160
161     // notify JTable that model has changed
162     fireTableStructureChanged();
163 } // end method setQuery
164
165 // close Statement and Connection
166 public void disconnectFromDatabase()
167 {
168     if ( connectedToDatabase )
169     {
170         // close Statement and Connection
171         try
172         {
173             resultSet.close();
174             statement.close();
175             connection.close();
176         } // end try
177         catch ( SQLException sqlException )
178         {
```

Fig. 28.25 | A TableModel1 that supplies ResultSet data to a JTable. (Part 4 of 5.)

```

179         sqlException.printStackTrace();
180     } // end catch
181     finally // update database connection status
182     {
183         connectedToDatabase = false;
184     } // end finally
185   } // end if
186 } // end method disconnectFromDatabase
187 } // end class ResultSetTableModel

```

Fig. 28.25 | A `TableModel` that supplies `ResultSet` data to a `JTable`. (Part 5 of 5.)

The `ResultSetTableModel` constructor (lines 30–46) accepts four `String` arguments—the URL of the database, the username, the password and the default query to perform. The constructor throws any exceptions that occur in its body back to the application that created the `ResultSetTableModel` object, so that the application can determine how to handle the exception (e.g., report an error and terminate the application). Line 34 establishes a connection to the database. Lines 37–39 invoke `Connection` method `createStatement` to create a `Statement` object. This example uses a version of method `createStatement` that takes two arguments—the result set type and the result set concurrency. The **result set type** (Fig. 28.26) specifies whether the `ResultSet`'s cursor is able to scroll in both directions or forward only and whether the `ResultSet` is sensitive to changes made to the underlying data.

ResultSet constant	Description
<code>TYPE_FORWARD_ONLY</code>	Specifies that a <code>ResultSet</code> 's cursor can move only in the forward direction (i.e., from the first to the last row in the <code>ResultSet</code>).
<code>TYPE_SCROLL_INSENSITIVE</code>	Specifies that a <code>ResultSet</code> 's cursor can scroll in either direction and that the changes made to the underlying data during <code>ResultSet</code> processing are not reflected in the <code>ResultSet</code> unless the program queries the database again.
<code>TYPE_SCROLL_SENSITIVE</code>	Specifies that a <code>ResultSet</code> 's cursor can scroll in either direction and that the changes made to the underlying data during <code>ResultSet</code> processing are reflected immediately in the <code>ResultSet</code> .

Fig. 28.26 | `ResultSet` constants for specifying `ResultSet` type.



Portability Tip 28.3

Some JDBC drivers do not support scrollable `ResultSet`s. In such cases, the driver typically returns a `ResultSet` in which the cursor can move only forward. For more information, see your database driver documentation.



Common Programming Error 28.11

Attempting to move the cursor backward through a `ResultSet` when the database driver does not support backward scrolling causes a `SQLException`.

`ResultSets` that are sensitive to changes reflect those changes immediately after they're made with methods of interface `ResultSet`. If a `ResultSet` is insensitive to changes, the

query that produced the `ResultSet` must be executed again to reflect any changes made. The **result set concurrency** (Fig. 28.27) specifies whether the `ResultSet` can be updated with `ResultSet`'s update methods.

ResultSet static concurrency constant	Description
<code>CONCUR_READ_ONLY</code>	Specifies that a <code>ResultSet</code> cannot be updated (i.e., changes to the <code>ResultSet</code> contents cannot be reflected in the database with <code>ResultSet</code> 's update methods).
<code>CONCUR_UPDATABLE</code>	Specifies that a <code>ResultSet</code> can be updated (i.e., changes to the <code>ResultSet</code> contents can be reflected in the database with <code>ResultSet</code> 's update methods).

Fig. 28.27 | `ResultSet` constants for specifying result properties.



Portability Tip 28.4

Some JDBC drivers do not support updatable ResultSets. In such cases, the driver typically returns a read-only ResultSet. For more information, see your database driver documentation.



Common Programming Error 28.12

Attempting to update a ResultSet when the database driver does not support updatable ResultSets causes SQLFeatureNotSupportedException.

This example uses a `ResultSet` that is scrollable, insensitive to changes and read only. Line 45 invokes our method `setQuery` (lines 144–163) to perform the default query.

Method `getColumnClass` (lines 49–69) returns a `Class` object that represents the superclass of all objects in a particular column. The `JTable` uses this information to configure the default cell renderer and cell editor for that column in the `JTable`. Line 58 uses `ResultSetMetaData` method `getColumnName` to obtain the fully qualified class name for the specified column. Line 61 loads the class and returns the corresponding `Class` object. If an exception occurs, the catch in lines 63–66 prints a stack trace and line 68 returns `Object.class`—the `Class` instance that represents class `Object`—as the default type. [Note: Line 58 uses the argument `column + 1`. Like arrays, `JTable` row and column numbers are counted from 0. However, `ResultSet` row and column numbers are counted from 1. Thus, when processing `ResultSet` rows or columns for use in a `JTable`, it is necessary to add 1 to the row or column number to manipulate the appropriate `ResultSet` row or column.]

Method `getColumnCount` (lines 72–89) returns the number of columns in the model's underlying `ResultSet`. Line 81 uses `ResultSetMetaData` method `getColumnCount` to obtain the number of columns in the `ResultSet`. If an exception occurs, the catch in lines 83–86 prints a stack trace and line 88 returns 0 as the default number of columns.

Method `getColumnName` (lines 92–109) returns the name of the column in the model's underlying `ResultSet`. Line 101 uses `ResultSetMetaData` method `getColumnName` to obtain the column name from the `ResultSet`. If an exception occurs, the catch in lines 103–106 prints a stack trace and line 108 returns the empty string as the default column name.

Method `getRowCount` (lines 112–119) returns the number of rows in the model's underlying `ResultSet`. When method `setQuery` (lines 144–163) performs a query, it stores the number of rows in variable `numberOfRows`.

Method `getValueAt` (lines 122–141) returns the `Object` in a particular row and column of the model's underlying `ResultSet`. Line 132 uses `ResultSet` method `absolute` to position the `ResultSet` cursor at a specific row. Line 133 uses `ResultSet` method `getObject` to obtain the `Object` in a specific column of the current row. If an exception occurs, the catch in lines 135–138 prints a stack trace and line 140 returns an empty string as the default value.

Method `setQuery` (lines 144–163) executes the query it receives as an argument to obtain a new `ResultSet` (line 152). Line 155 gets the `ResultSetMetaData` for the new `ResultSet`. Line 158 uses `ResultSet` method `last` to position the `ResultSet` cursor at the last row in the `ResultSet`. [Note: This can be slow if the table contains many rows.] Line 159 uses `ResultSet` method `getRow` to obtain the row number for the current row in the `ResultSet`. Line 162 invokes method `fireTableStructureChanged` (inherited from class `AbstractTableModel`) to notify any `JTable` using this `ResultSetTableModel` object as its model that the structure of the model has changed. This causes the `JTable` to repopulate its rows and columns with the new `ResultSet` data. Method `setQuery` throws any exceptions that occur in its body back to the application that invoked `setQuery`.

Method `disconnectFromDatabase` (lines 166–186) implements an appropriate termination method for class `ResultSetTableModel`. A class designer should provide a `public` method that clients of the class must invoke explicitly to free resources that an object has used. In this case, method `disconnectFromDatabase` closes the `ResultSet`, `Statement` and `Connection` (lines 173–175), which are considered limited resources. Clients of the `ResultSetTableModel` class should always invoke this method when the instance of this class is no longer needed. Before releasing resources, line 168 verifies whether the connection is already terminated. If not, the method proceeds. Note that the other methods in class `ResultSetTableModel` each throw an `IllegalStateException` if `connectedToDatabase` is `false`. Method `disconnectFromDatabase` sets `connectedToDatabase` to `false` (line 183) to ensure that clients do not use an instance of `ResultSetTableModel` after that instance has already been terminated. `IllegalStateException` is an exception from the Java libraries that is appropriate for indicating this error condition.

DisplayQueryResults Class

Class `DisplayQueryResults` (Fig. 28.28) implements the application's GUI and interacts with the `ResultSetTableModel` via a `JTable` object. This application also demonstrates the `JTable` sorting and filtering capabilities introduced in Java SE 6.

```
1 // Fig. 28.28: DisplayQueryResults.java
2 // Display the contents of the Authors table in the books database.
3 import java.awt.BorderLayout;
4 import java.awt.event.ActionListener;
5 import java.awt.event.ActionEvent;
6 import java.awt.event.WindowAdapter;
7 import java.awt.event.WindowEvent;
8 import java.sql.SQLException;
9 import java.util.regex.PatternSyntaxException;
```

Fig. 28.28 | Displays contents of the database books. (Part I of 6.)

```
10 import javax.swing.JFrame;
11 import javax.swing.JTextArea;
12 import javax.swing.JScrollPane;
13 import javax.swing.ScrollPaneConstants;
14 import javax.swing.JTable;
15 import javax.swing.JOptionPane;
16 import javax.swing.JButton;
17 import javax.swing.Box;
18 import javax.swing.JLabel;
19 import javax.swing.JTextField;
20 import javax.swing.RowFilter;
21 import javax.swing.table.TableRowSorter;
22 import javax.swing.table.TableModel;
23
24 public class DisplayQueryResults extends JFrame
25 {
26     // database URL, username and password
27     static final String DATABASE_URL = "jdbc:mysql://localhost/books";
28     static final String USERNAME = "deitel";
29     static final String PASSWORD = "deitel";
30
31     // default query retrieves all data from Authors table
32     static final String DEFAULT_QUERY = "SELECT * FROM Authors";
33
34     private ResultSetTableModel tableModel;
35     private JTextArea queryArea;
36
37     // create ResultSetTableModel and GUI
38     public DisplayQueryResults()
39     {
40         super( "Displaying Query Results" );
41
42         // create ResultSetTableModel and display database table
43         try
44         {
45             // create TableModel for results of query SELECT * FROM Authors
46             tableModel = new ResultSetTableModel( DATABASE_URL,
47                     USERNAME, PASSWORD, DEFAULT_QUERY );
48
49             // set up JTextArea in which user types queries
50             queryArea = new JTextArea( DEFAULT_QUERY, 3, 100 );
51             queryArea.setWrapStyleWord( true );
52             queryArea.setLineWrap( true );
53
54             JScrollPane scrollPane = new JScrollPane( queryArea,
55                     ScrollPaneConstants.VERTICAL_SCROLLBAR_AS_NEEDED,
56                     ScrollPaneConstants.HORIZONTAL_SCROLLBAR_NEVER );
57
58             // set up JButton for submitting queries
59             JButton submitButton = new JButton( "Submit Query" );
60
61             // create Box to manage placement of queryArea and
62             // submitButton in GUI
63             Box boxNorth = Box.createHorizontalBox();
```

Fig. 28.28 | Displays contents of the database books. (Part 2 of 6.)

```
64      boxNorth.add( scrollPane );
65      boxNorth.add( submitButton );
66
67      // create JTable based on the tableModel
68      JTable resultTable = new JTable( tableModel );
69
70      JLabel filterLabel = new JLabel( "Filter:" );
71      final JTextField filterText = new JTextField();
72      JButton filterButton = new JButton( "Apply Filter" );
73      Box boxSouth = Box.createHorizontalBox();
74
75      boxSouth.add( filterLabel );
76      boxSouth.add( filterText );
77      boxSouth.add( filterButton );
78
79      // place GUI components on content pane
80      add( boxNorth, BorderLayout.NORTH );
81      add( new JScrollPane( resultTable ), BorderLayout.CENTER );
82      add( boxSouth, BorderLayout.SOUTH );
83
84      // create event listener for submitButton
85      submitButton.addActionListener(
86
87          new ActionListener()
88          {
89              // pass query to table model
90              public void actionPerformed( ActionEvent event )
91              {
92                  // perform a new query
93                  try
94                  {
95                      tableModel.setQuery( queryArea.getText() );
96                  } // end try
97                  catch ( SQLException sqlException )
98                  {
99                      JOptionPane.showMessageDialog( null,
100                          sqlException.getMessage(), "Database error",
101                          JOptionPane.ERROR_MESSAGE );
102
103                      // try to recover from invalid user query
104                      // by executing default query
105                      try
106                      {
107                          tableModel.setQuery( DEFAULT_QUERY );
108                          queryArea.setText( DEFAULT_QUERY );
109                      } // end try
110                      catch ( SQLException sqlException2 )
111                      {
112                          JOptionPane.showMessageDialog( null,
113                              sqlException2.getMessage(), "Database error",
114                              JOptionPane.ERROR_MESSAGE );
115
116                      // ensure database connection is closed
117                      tableModel.disconnectFromDatabase();
```

Fig. 28.28 | Displays contents of the database books. (Part 3 of 6.)

```
118                     System.exit( 1 ); // terminate application
119             } // end inner catch
120         } // end outer catch
121     } // end actionPerformed
122 } // end ActionListener inner class
123 ); // end call to addActionListener
124
125
126     final TableRowSorter<TableModel> sorter =
127         new TableRowSorter<TableModel>( tableView );
128     resultTable.setRowSorter( sorter );
129     setSize( 500, 250 ); // set window size
130     setVisible( true ); // display window
131
132     // create listener for filterButton
133     filterButton.addActionListener(
134         new ActionListener()
135     {
136         // pass filter text to listener
137         public void actionPerformed( ActionEvent e )
138     {
139         String text = filterText.getText();
140
141         if ( text.length() == 0 )
142             sorter.setRowFilter( null );
143         else
144         {
145             try
146             {
147                 sorter.setRowFilter(
148                     RowFilter.regexFilter( text ) );
149             } // end try
150             catch ( PatternSyntaxException pse )
151             {
152                 JOptionPane.showMessageDialog( null,
153                     "Bad regex pattern", "Bad regex pattern",
154                     JOptionPane.ERROR_MESSAGE );
155             } // end catch
156         } // end else
157     } // end method actionPerformed
158 } // end anonymous inner class
159 ); // end call to addActionListener
160 } // end try
161 catch ( SQLException sqlException )
162 {
163     JOptionPane.showMessageDialog( null, sqlException.getMessage(),
164         "Database error", JOptionPane.ERROR_MESSAGE );
165
166     // ensure database connection is closed
167     tableView.disconnectFromDatabase();
168
169     System.exit( 1 ); // terminate application
170 } // end catch
```

Fig. 28.28 | Displays contents of the database books. (Part 4 of 6.)

```

172     // dispose of window when user quits application (this overrides
173     // the default of HIDE_ON_CLOSE)
174     setDefaultCloseOperation( DISPOSE_ON_CLOSE );
175
176     // ensure database connection is closed when user quits application
177     addWindowListener(
178
179         new WindowAdapter()
180     {
181         // disconnect from database and exit when window has closed
182         public void windowClosed( WindowEvent event )
183         {
184             tableModel.disconnectFromDatabase();
185             System.exit( 0 );
186         } // end method windowClosed
187     } // end WindowAdapter inner class
188 ); // end call to addWindowListener
189 } // end DisplayQueryResults constructor
190
191 // execute application
192 public static void main( String args[] )
193 {
194     new DisplayQueryResults();
195 } // end main
196 } // end class DisplayQueryResults

```

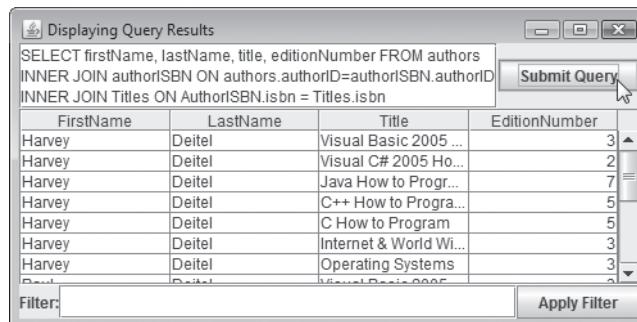
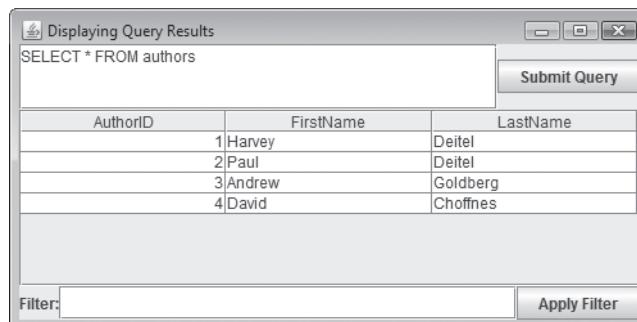


Fig. 28.28 | Displays contents of the database books. (Part 5 of 6.)

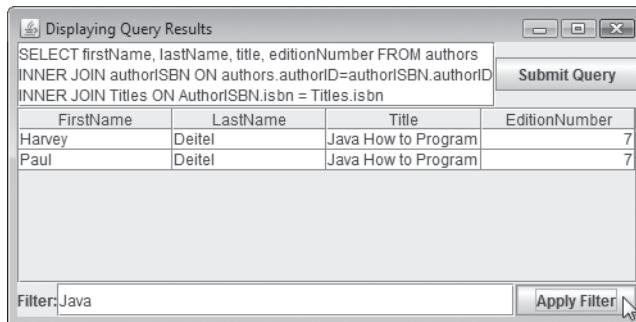


Fig. 28.28 | Displays contents of the database books. (Part 6 of 6.)

Lines 27–29 and 32 declare the URL, username, password and default query that are passed to the `ResultSetTableModel` constructor to make the initial connection to the database and perform the default query. The `DisplayQueryResults` constructor (lines 38–189) creates a `ResultSetTableModel` object and the GUI for the application. Line 68 creates the `JTable` object and passes a `ResultSetTableModel` object to the `JTable` constructor, which then registers the `JTable` as a listener for `TableModelEvents` generated by the `ResultSetTableModel`.

Note that the local variables `filterText` (line 71) and `sorter` (lines 126–127) are declared `final`. These are both used from an event handler that is implemented as an anonymous inner class (lines 134–158). Any local variable that will be used in an anonymous inner class *must* be declared `final`; otherwise, a compilation error occurs.

Lines 85–124 register an event handler for the `submitButton` that the user clicks to submit a query to the database. When the user clicks the button, method `actionPerformed` (lines 90–122) invokes method `setQuery` from the class `ResultSetTableModel` to execute the new query (line 95). If the user's query fails (e.g., because of a syntax error in the user's input), lines 107–108 execute the default query. If the default query also fails, there could be a more serious error, so line 117 ensures that the database connection is closed and line 119 exits the program. The screen captures in Fig. 28.28 show the results of two queries. The first screen capture shows the default query that retrieves all the data from table `Authors` of database `books`. The second screen capture shows a query that selects each author's first name and last name from the `Authors` table and combines that information with the title and edition number from the `Titles` table. Try entering your own queries in the text area and clicking the `Submit Query` button to execute the query.

As of Java SE 6, `JTables` now allow users to sort rows by the data in a specific column. Lines 126–127 use the `TableRowSorter` class (from package `javax.swing.table`) to create an object that uses our `ResultSetTableModel` to sort rows in the `JTable` that displays query results. When the user clicks the title of a particular `JTable` column, the `TableRowSorter` interacts with the underlying `TableModel` to reorder the rows based on the data in that column. Line 128 uses `JTable` method `setRowSorter` to specify the `TableRowSorter` for `resultTable`.

`JTables` can now show subsets of the data from the underlying `TableModel`. This is known as filtering the data. Lines 133–159 register an event handler for the `filterButton` that the user clicks to filter the data. In method `actionPerformed` (lines 137–157), line

139 obtains the filter text. If the user did not specify filter text, line 142 uses `JTable` method `setRowFilter` to remove any prior filter by setting the filter to `null`. Otherwise, lines 147–148 use `setRowFilter` to specify a `RowFilter` (from package `javax.swing`) based on the user's input. Class `RowFilter` provides several methods for creating filters. The static method `regexFilter` receives a `String` containing a regular expression pattern as its argument and an optional set of indices that specify which columns to filter. If no indices are specified, then all the columns are searched. In this example, the regular expression pattern is the text the user typed. Once the filter is set, the data displayed in the `JTable` is updated based on the filtered `TableModel`.

Lines 177–188 register a `WindowListener` for the `windowClosed` event, which occurs when the user closes the window. Since `WindowListeners` can handle several window events, we extend class `WindowAdapter` and override only the `windowClosed` event handler.

28.9 RowSet Interface

In the previous examples, you learned how to query a database by explicitly establishing a `Connection` to the database, preparing a `Statement` for querying the database and executing the query. In this section, we demonstrate the `RowSet interface`, which configures the database connection and prepares query statements automatically. The interface `RowSet` provides several *set* methods that allow you to specify the properties needed to establish a connection (such as the database URL, user name and password of the database) and create a `Statement` (such as a query). `RowSet` also provides several *get* methods that return these properties.

There are two types of `RowSet` objects—connected and disconnected. A `connected RowSet` object connects to the database once and remains connected while the object is in use. A `disconnected RowSet` object connects to the database, executes a query to retrieve the data from the database and then closes the connection. A program may change the data in a disconnected `RowSet` while it is disconnected. Modified data can be updated in the database after a disconnected `RowSet` reestablishes the connection with the database.

Package `javax.sql.rowset` contains two subinterfaces of `RowSet`—`JdbcRowSet` and `CachedRowSet`. `JdbcRowSet`, a connected `RowSet`, acts as a wrapper around a `ResultSet` object and allows you to scroll through and update the rows in the `ResultSet`. Recall that by default, a `ResultSet` object is nonscrollable and read only—you must explicitly set the result set type constant to `TYPE_SCROLL_INSENSITIVE` and set the result set concurrency constant to `CONCUR_UPDATABLE` to make a `ResultSet` object scrollable and updatable. A `JdbcRowSet` object is scrollable and updatable by default. `CachedRowSet`, a disconnected `RowSet`, caches the data of a `ResultSet` in memory and disconnects from the database. Like `JdbcRowSet`, a `CachedRowSet` object is scrollable and updatable by default. A `CachedRowSet` object is also serializable, so it can be passed between Java applications through a network, such as the Internet. However, `CachedRowSet` has a limitation—the amount of data that can be stored in memory is limited. Package `javax.sql.rowset` contains three other subinterfaces of `RowSet`. For details of these interfaces, visit java.sun.com/javase/6/docs/technotes/guides/jdbc/getstart/rowsetImpl.html.



Portability Tip 28.5

A `RowSet` can provide scrolling capability for drivers that do not support scrollable `ResultSets`.

Figure 28.29 reimplements the example of Fig. 28.23 using a `RowSet`. Rather than establish the connection and create a `Statement` explicitly, Fig. 28.29 uses a `JdbcRowSet` object to create a `Connection` and a `Statement` automatically.

```

1 // Fig. 28.29: JdbcRowSetTest.java
2 // Displaying the contents of the Authors table using JdbcRowSet.
3 import java.sql.ResultSetMetaData;
4 import java.sql.SQLException;
5 import javax.sql.rowset.JdbcRowSet;
6 import com.sun.rowset.JdbcRowSetImpl; // Sun's JdbcRowSet implementation
7
8 public class JdbcRowSetTest
9 {
10    // JDBC driver name and database URL
11    static final String DATABASE_URL = "jdbc:mysql://localhost/books";
12    static final String USERNAME = "deitel";
13    static final String PASSWORD = "deitel";
14
15    // constructor connects to database, queries database, processes
16    // results and displays results in window
17    public JdbcRowSetTest()
18    {
19        // connect to database books and query database
20        try
21        {
22            // specify properties of JdbcRowSet
23            JdbcRowSet rowSet = new JdbcRowSetImpl();
24            rowSet.setUrl( DATABASE_URL ); // set database URL
25            rowSet.setUsername( USERNAME ); // set username
26            rowSet.setPassword( PASSWORD ); // set password
27            rowSet.setCommand( "SELECT * FROM Authors" ); // set query
28            rowSet.execute(); // execute query
29
30            // process query results
31            ResultSetMetaData metaData = rowSet.getMetaData();
32            int numberOfColumns = metaData.getColumnCount();
33            System.out.println( "Authors Table of Books Database:\n" );
34
35            // display rowset header
36            for ( int i = 1; i <= numberOfColumns; i++ )
37                System.out.printf( "%-8s\t", metaData.getColumnName( i ) );
38            System.out.println();
39
40            // display each row
41            while ( rowSet.next() )
42            {
43                for ( int i = 1; i <= numberOfColumns; i++ )
44                    System.out.printf( "%-8s\t", rowSet.getObject( i ) );
45                System.out.println();
46            } // end while
47
48            // close the underlying ResultSet, Statement and Connection
49            rowSet.close();
50        } // end try

```

Fig. 28.29 | Displaying the Authors table using `JdbcRowSet`. (Part I of 2.)

```

51     catch ( SQLException sqlException )
52     {
53         sqlException.printStackTrace();
54         System.exit( 1 );
55     } // end catch
56 } // end DisplayAuthors constructor
57
58 // launch the application
59 public static void main( String args[] )
60 {
61     JdbcRowSetTest application = new JdbcRowSetTest();
62 } // end main
63 } // end class JdbcRowSetTest

```

Authors Table of Books Database:

AuthorID	FirstName	LastName
1	Harvey	Deitel
2	Paul	Deitel
3	Andrew	Goldberg
4	David	Choffnes

Fig. 28.29 | Displaying the Authors table using `JdbcRowSet`. (Part 2 of 2.)

The package `com.sun.rowset` provides Sun's reference implementations of the interfaces in package `javax.sql.rowset`. Line 23 uses Sun's reference implementation of the `JdbcRowSet` interface—`JdbcRowSetImpl`—to create a `JdbcRowSet` object. We used class `JdbcRowSetImpl` here to demonstrate the capability of the `JdbcRowSet` interface. Other databases may provide their own RowSet implementations.

Lines 24–26 set the RowSet properties that the `DriverManager` uses to establish a database connection. Line 24 invokes `JdbcRowSet` method `setUrl` to specify the database URL. Line 25 invokes `JdbcRowSet` method `setUsername` to specify the username. Line 26 invokes `JdbcRowSet` method `setPassword` to specify the password. Line 27 invokes `JdbcRowSet` method `setCommand` to specify the SQL query that will be used to populate the RowSet. Line 28 invokes `JdbcRowSet` method `execute` to execute the SQL query. Method `execute` performs four actions—it establishes a `Connection` to the database, prepares the `Query Statement`, executes the query and stores the `ResultSet` returned by query. The `Connection`, `Statement` and `ResultSet` are encapsulated in the `JdbcRowSet` object.

The remaining code is almost identical to Fig. 28.23, except that line 31 obtains a `ResultSetMetaData` object from the `JdbcRowSet`, line 41 uses the `JdbcRowSet`'s `next` method to get the next row of the result and line 44 uses the `JdbcRowSet`'s `getObject` method to obtain a column's value. Line 49 invokes `JdbcRowSet` method `close`, which closes the RowSet's encapsulated `ResultSet`, `Statement` and `Connection`. In a `CachedRowSet`, invoking `close` also releases the resources held by that RowSet. Note that the output of this application is the same as that of Fig. 28.23.

28.10 Java DB/Apache Derby

As of JDK 6, Sun Microsystems now bundles the open-source, pure Java database **Java DB** (the Sun branded version of Apache Derby) with the JDK. In Section 28.11, we use Java

DB to demonstrate `PreparedStatement`s. Before you can execute the application in the next section, you must set up the AddressBook database in Java DB. Section 28.11 uses the embedded version of Java DB. There is also a network version that executes similarly to the MySQL DBMS introduced earlier in the chapter. For the purpose of the following steps, we assume you are running Microsoft Windows with Java installed in its default location. [Note: Mac OS X users may need to download and install Apache Derby (db.apache.org/derby/) if Java DB is not installed by default. Linux users should visit developers.sun.com/javadb/reference/docs/ and read “Getting Started with Java DB” for installation and configuration instructions.]

1. Java DB comes with several batch files to configure and run it. Before executing these batch files from a command prompt, you must set the environment variable `JAVA_HOME` to refer to the JDK’s `C:\Program Files\Java\jdk1.6.0_11` installation directory. [Note: You may need to change this based on the version of the JDK installed on your computer.]
2. Open the batch file `setEmbeddedCP.bat` (located in `C:\Program Files\Sun\JavaDB\bin`) in a text editor such as Notepad. Locate the line

```
@rem set DERBY_INSTALL=
```

and change it to

```
@set DERBY_INSTALL=C:\Program Files\Sun\JavaDB
```

Save your changes and close this file.

3. Open a Command Prompt and change directories to `C:\Program Files\Sun\JavaDB\bin`. Then, type `setEmbeddedCP.bat` and press *Enter* to set the environment variables required by Java DB.
4. An embedded Java DB database must reside in the same location as the application that manipulates the database. For this reason, change to the directory that contains the code for Figs. 28.30–28.32. This directory contains a SQL script `address.sql` that builds the AddressBook database.
5. Execute the command

```
"C:\Program Files\Sun\JavaDB\bin\ij"
```

to start the command-line tool for interacting with Java DB. The double quotes are necessary because the path contains a space. This will display the `ij>` prompt.

6. At the `ij>` prompt type
- ```
connect 'jdbc:derby:AddressBook;create=true;user=deitel';
password=deitel';
```
- to create the AddressBook database in the current directory. This command also creates the user `deitel` with the password `deitel` for accessing the database.

  7. To create the database table and insert sample data in the database, type
- ```
run 'address.sql';
```
8. To terminate the Java DB command-line tool, type
- ```
exit;
```
- © 2010 Pearson Education, Inc., Upper Saddle River, NJ. All Rights Reserved.

You are now ready to execute the AddressBook application in Section 28.11. Note that MySQL (or any other database that supports JDBC `PreparedStatement`) could also be used in Section 28.11.

## 28.11 PreparedStatements

Interface `PreparedStatement` enables you to create compiled SQL statements that execute more efficiently than `Statement` objects. `PreparedStatement`s also can specify parameters, making them more flexible than `Statements`. Programs can execute the same query repeatedly with different parameter values. For example, in the books database, you might want to locate all book titles for an author with a specific last name and first name, and you might want to execute that query for several authors. With a `PreparedStatement`, that query is defined as follows:

```
PreparedStatement authorBooks = connection.prepareStatement(
 "SELECT LastName, FirstName, Title " +
 "FROM Authors INNER JOIN AuthorISBN " +
 "ON Authors.AuthorID=AuthorISBN.AuthorID " +
 "INNER JOIN Titles " +
 "ON AuthorISBN.ISBN=Titles.ISBN " +
 "WHERE LastName = ? AND FirstName = ?");
```

The two question marks (?) in the preceding SQL statement's last line are placeholders for values that will be passed as part of the query to the database. Before executing a `PreparedStatement`, the program must specify the parameter values by using the `PreparedStatement` interface's *set* methods.

For the preceding query, both parameters are strings that can be set with `PreparedStatement` method `setString` as follows:

```
authorBooks.setString(1, "Deitel");
authorBooks.setString(2, "Paul");
```

Method `setString`'s first argument represents the parameter number being set, and the second argument is that parameter's value. Parameter numbers are counted from 1, starting with the first question mark (?). When the program executes the preceding `PreparedStatement` with the parameter values shown here, the SQL passed to the database is

```
SELECT LastName, FirstName, Title
FROM Authors INNER JOIN AuthorISBN
 ON Authors.AuthorID=AuthorISBN.AuthorID
INNER JOIN Titles
 ON AuthorISBN.ISBN=Titles.ISBN
WHERE LastName = 'Deitel' AND FirstName = 'Paul'
```

Method `setString` automatically escapes String parameter values as necessary. For example, if the last name is O'Brien, the statement

```
authorBooks.setString(1, "O'Brien");
```

escapes the ' character in O'Brien by replacing it with two single-quote characters.



### Performance Tip 28.2

*PreparedStatement are more efficient than Statements when executing SQL statements multiple times and with different parameter values.*



### Error-Prevention Tip 28.2

*Use PreparedStatements with parameters for queries that receive String values as arguments to ensure that the Strings are quoted properly in the SQL statement.*

Interface `PreparedStatement` provides `set` methods for each supported SQL type. It is important to use the `set` method that is appropriate for the parameter's SQL type in the database—`SQLExceptions` occur when a program attempts to convert a parameter value to an incorrect type. For a complete list of interface `PreparedStatement`'s `set` methods, see [java.sun.com/javase/6/docs/api/java/sql/PreparedStatement.html](http://java.sun.com/javase/6/docs/api/java/sql/PreparedStatement.html).

#### *Address Book Application that Uses `PreparedStatement`*

We now present an address book application that enables you to browse existing entries, add new entries and search for entries with a specific last name. Our `AddressBook` Java DB database contains an `Addresses` table with the columns `addressID`, `FirstName`, `LastName`, `Email` and `PhoneNumber`. The column `addressID` is a so-called identity column. This is the SQL standard way to represent an autoincremented column. The SQL script we provide for this database uses the SQL `IDENTITY` keyword to mark the `addressID` column as an identity column. For more information on using the `IDENTITY` keyword and creating databases, see the Java DB Developer's Guide at [developers.sun.com/docs/javadb/10.4.2.0/devguide/derbydev.pdf](http://developers.sun.com/docs/javadb/10.4.2.0/devguide/derbydev.pdf).

Our address book application consists of three classes—`Person` (Fig. 28.30), `PersonQueries` (Fig. 28.31) and `AddressBookDisplay` (Fig. 28.32). Class `Person` is a simple class that represents one person in the address book. The class contains fields for the address ID, first name, last name, email address and phone number, as well as `set` and `get` methods for manipulating these fields.

---

```

1 // Fig. 28.30: Person.java
2 // Person class that represents an entry in an address book.
3 public class Person
4 {
5 private int addressID;
6 private String firstName;
7 private String lastName;
8 private String email;
9 private String phoneNumber;
10
11 // no-argument constructor
12 public Person()
13 {
14 // end no-argument Person constructor
15
16 // constructor
17 public Person(int id, String first, String last,
18 String emailAddress, String phone)
19 {
20 setAddressID(id);
21 setFirstName(first);
22 setLastName(last);
23 setEmail(emailAddress);

```

**Fig. 28.30** | Person class that represents an entry in an AddressBook. (Part 1 of 3.)

```
24 setPhoneNumber(phone);
25 } // end five-argument Person constructor
26
27 // sets the addressID
28 public void setAddressID(int id)
29 {
30 addressID = id;
31 } // end method setAddressID
32
33 // returns the addressID
34 public int getAddressID()
35 {
36 return addressID;
37 } // end method getAddressID
38
39 // sets the firstName
40 public void setFirstName(String first)
41 {
42 firstName = first;
43 } // end method setFirstName
44
45 // returns the first name
46 public String getFirstName()
47 {
48 return firstName;
49 } // end method getFirstName
50
51 // sets the lastName
52 public void setLastName(String last)
53 {
54 lastName = last;
55 } // end method setLastName
56
57 // returns the last name
58 public String getLastName()
59 {
60 return lastName;
61 } // end method getLastName
62
63 // sets the email address
64 public void setEmail(String emailAddress)
65 {
66 email = emailAddress;
67 } // end method setEmail
68
69 // returns the email address
70 public String getEmail()
71 {
72 return email;
73 } // end method getEmail
74
75 // sets the phone number
76 public void setPhoneNumber(String phone)
77 {
```

**Fig. 28.30** | Person class that represents an entry in an AddressBook. (Part 2 of 3.)

---

```

78 phoneNumber = phone;
79 } // end method setPhoneNumber
80
81 // returns the phone number
82 public String getPhoneNumber()
83 {
84 return phoneNumber;
85 } // end method getPhoneNumber
86 } // end class Person

```

---

**Fig. 28.30** | Person class that represents an entry in an AddressBook. (Part 3 of 3.)***Class PersonQueries***

Class `PersonQueries` (Fig. 28.31) manages the address book application's database connection and creates the `PreparedStatement`s that the application uses to interact with the database. Lines 18–20 declare three `PreparedStatement` variables. The constructor (lines 23–49) connects to the database at lines 27–28.

Lines 31–32 invoke `Connection` method `prepareStatement` to create the `PreparedStatement` named `selectAllPeople` that selects all the rows in the `Addresses` table. Lines 35–36 create the `PreparedStatement` named `selectPeopleByLastName` with a parameter. This statement selects all the rows in the `Addresses` table that match a particular last name. Notice the `?` character that is used to specify the last-name parameter. Lines 39–42 create the `PreparedStatement` named `insertNewPerson` with four parameters that represent the first name, last name, email address and phone number for a new entry. Again, notice the `?` characters used to represent these parameters.

---

```

1 // Fig. 28.31: PersonQueries.java
2 // PreparedStatements used by the Address Book application.
3 import java.sql.Connection;
4 import java.sql.DriverManager;
5 import java.sql.PreparedStatement;
6 import java.sql.ResultSet;
7 import java.sql.SQLException;
8 import java.util.List;
9 import java.util.ArrayList;
10
11 public class PersonQueries
12 {
13 private static final String URL = "jdbc:derby:AddressBook";
14 private static final String USERNAME = "deitel";
15 private static final String PASSWORD = "deitel";
16
17 private Connection connection = null; // manages connection
18 private PreparedStatement selectAllPeople = null;
19 private PreparedStatement selectPeopleByLastName = null;
20 private PreparedStatement insertNewPerson = null;
21
22 // constructor
23 public PersonQueries()
24 {

```

---

**Fig. 28.31** | PreparedStatements used by the Address Book application. (Part 1 of 4.)

```
25 try
26 {
27 connection =
28 DriverManager.getConnection(URL, USERNAME, PASSWORD);
29
30 // create query that selects all entries in the AddressBook
31 selectAllPeople =
32 connection.prepareStatement("SELECT * FROM Addresses");
33
34 // create query that selects entries with a specific last name
35 selectPeopleByLastName = connection.prepareStatement(
36 "SELECT * FROM Addresses WHERE LastName = ?");
37
38 // create insert that adds a new entry into the database
39 insertNewPerson = connection.prepareStatement(
40 "INSERT INTO Addresses " +
41 "(FirstName, LastName, Email, PhoneNumber) " +
42 "VALUES (?, ?, ?, ?)");
43 } // end try
44 catch (SQLException sqlException)
45 {
46 sqlException.printStackTrace();
47 System.exit(1);
48 } // end catch
49 } // end PersonQueries constructor
50
51 // select all of the addresses in the database
52 public List< Person > getAllPeople()
53 {
54 List< Person > results = null;
55 ResultSet resultSet = null;
56
57 try
58 {
59 // executeQuery returns ResultSet containing matching entries
60 resultSet = selectAllPeople.executeQuery();
61 results = new ArrayList< Person >();
62
63 while (resultSet.next())
64 {
65 results.add(new Person(
66 resultSet.getInt("addressID"),
67 resultSet.getString("FirstName"),
68 resultSet.getString("LastName"),
69 resultSet.getString("Email"),
70 resultSet.getString("PhoneNumber")));
71 } // end while
72 } // end try
73 catch (SQLException sqlException)
74 {
75 sqlException.printStackTrace();
76 } // end catch
77 finally
78 {
```

---

**Fig. 28.31 |** PreparedStatements used by the Address Book application. (Part 2 of 4.)

```
79 try
80 {
81 resultSet.close();
82 } // end try
83 catch (SQLException sqlException)
84 {
85 sqlException.printStackTrace();
86 close();
87 } // end catch
88 } // end finally
89
90 return results;
91 } // end method getAllPeople
92
93 // select person by last name
94 public List< Person > getPeopleByLastName(String name)
95 {
96 List< Person > results = null;
97 ResultSet resultSet = null;
98
99 try
100 {
101 selectPeopleByLastName.setString(1, name); // specify last name
102
103 // executeQuery returns ResultSet containing matching entries
104 resultSet = selectPeopleByLastName.executeQuery();
105
106 results = new ArrayList< Person >();
107
108 while (resultSet.next())
109 {
110 results.add(new Person(resultSet.getInt("addressID"),
111 resultSet.getString("FirstName"),
112 resultSet.getString("LastName"),
113 resultSet.getString("Email"),
114 resultSet.getString("PhoneNumber")));
115 } // end while
116 } // end try
117 catch (SQLException sqlException)
118 {
119 sqlException.printStackTrace();
120 } // end catch
121 finally
122 {
123 try
124 {
125 resultSet.close();
126 } // end try
127 catch (SQLException sqlException)
128 {
129 sqlException.printStackTrace();
130 close();
131 } // end catch
132 } // end finally
```

---

**Fig. 28.31 |** PreparedStatements used by the Address Book application. (Part 3 of 4.)

```

133 return results;
134 } // end method getPeopleByName
135
136
137 // add an entry
138 public int addPerson(
139 String fname, String lname, String email, String num)
140 {
141 int result = 0;
142
143 // set parameters, then execute insertNewPerson
144 try
145 {
146 insertNewPerson.setString(1, fname);
147 insertNewPerson.setString(2, lname);
148 insertNewPerson.setString(3, email);
149 insertNewPerson.setString(4, num);
150
151 // insert the new entry; returns # of rows updated
152 result = insertNewPerson.executeUpdate();
153 } // end try
154 catch (SQLException sqlException)
155 {
156 sqlException.printStackTrace();
157 close();
158 } // end catch
159
160 return result;
161 } // end method addPerson
162
163 // close the database connection
164 public void close()
165 {
166 try
167 {
168 connection.close();
169 } // end try
170 catch (SQLException sqlException)
171 {
172 sqlException.printStackTrace();
173 } // end catch
174 } // end method close
175 } // end class PersonQueries

```

**Fig. 28.31** | PreparedStatements used by the Address Book application. (Part 4 of 4.)

Method `getAllPeople` (lines 52–91) executes `PreparedStatement selectAllPeople` (line 60) by calling method `executeQuery`, which returns a `ResultSet` containing the rows that match the query (in this case, all the rows in the `Addresses` table). Lines 61–71 place the query results in an `ArrayList` of `Person` objects, which is returned to the caller at line 90. Method `getPeopleByLastName` (lines 94–135) uses `PreparedStatement` method `setString` to set the parameter to `selectPeopleByLastName` (line 101). Then, line 104 executes the query and lines 106–115 place the query results in an `ArrayList` of `Person` objects. Line 134 returns the `ArrayList` to the caller.

Method `addPerson` (lines 138–161) uses `PreparedStatement` method `setString` (lines 146–149) to set the parameters for the `insertNewPerson PreparedStatement`. Line 152 uses `PreparedStatement` method `executeUpdate` to insert the new record. This method returns an integer indicating the number of rows that were updated (or inserted) in the database. Method `close` (lines 164–174) simply closes the database connection.

### **Class AddressBookDisplay**

The `AddressBookDisplay` (Fig. 28.32) application uses a `PersonQueries` object to interact with the database. Line 59 creates the `PersonQueries` object. When the user presses the `Browse All Entries` JButton, the `browseButtonActionPerformed` handler (lines 309–335) is called. Line 313 calls the method `getAllPeople` on the `PersonQueries` object to obtain all the entries in the database. The user can then scroll through the entries using the `Previous` and `Next` JButtons. When the user presses the `Find` JButton, the `queryButtonActionPerformed` handler (lines 265–287) is called. Lines 267–268 call method `getPeopleByLastName` on the `PersonQueries` object to obtain the entries in the database that match the specified last name. If there are several such entries, the user can then scroll through them using the `Previous` and `Next` JButtons.

---

```
1 // Fig. 28.32: AddressBookDisplay.java
2 // A simple address book
3 import java.awt.event.ActionEvent;
4 import java.awt.event.ActionListener;
5 import java.awt.event.WindowAdapter;
6 import java.awt.event.WindowEvent;
7 import java.awt.FlowLayout;
8 import java.awt.GridLayout;
9 import java.util.List;
10 import javax.swing.JButton;
11 import javax.swing.Box;
12 import javax.swing.JFrame;
13 import javax.swing.JLabel;
14 import javax.swing.JPanel;
15 import javax.swing.JTextField;
16 import javax.swing.WindowConstants;
17 import javax.swingBoxLayout;
18 import javax.swing.BorderFactory;
19 import javax.swing.JOptionPane;
20
21 public class AddressBookDisplay extends JFrame
22 {
23 private Person currentEntry;
24 private PersonQueries personQueries;
25 private List< Person > results;
26 private int numberOfEntries = 0;
27 private int currentEntryIndex;
28
29 private JButton browseButton;
30 private JLabel emailLabel;
31 private JTextField emailTextField;
32 private JLabel firstNameLabel;
33 private JTextField firstNameTextField;
```

**Fig. 28.32** | A simple address book. (Part 1 of 8.)

```
34 private JLabel idLabel;
35 private JTextField idTextField;
36 private JTextField indexTextField;
37 private JLabel lastNameLabel;
38 private JTextField lastNameTextField;
39 private JTextField maxTextField;
40 private JButton nextButton;
41 private JLabel ofLabel;
42 private JLabel phoneLabel;
43 private JTextField phoneTextField;
44 private JButton previousButton;
45 private JButton queryButton;
46 private JLabel queryLabel;
47 private JPanel queryPanel;
48 private JPanel navigatePanel;
49 private JPanel displayPanel;
50 private JTextField queryTextField;
51 private JButton insertButton;
52
53 // no-argument constructor
54 public AddressBookDisplay()
55 {
56 super("Address Book");
57
58 // establish database connection and set up PreparedStatements
59 personQueries = new PersonQueries();
60
61 // create GUI
62 navigatePanel = new JPanel();
63 previousButton = new JButton();
64 indexTextField = new JTextField(2);
65 ofLabel = new JLabel();
66 maxTextField = new JTextField(2);
67 nextButton = new JButton();
68 displayPanel = new JPanel();
69 idLabel = new JLabel();
70 idTextField = new JTextField(10);
71 firstNameLabel = new JLabel();
72 firstNameTextField = new JTextField(10);
73 lastNameLabel = new JLabel();
74 lastNameTextField = new JTextField(10);
75 emailLabel = new JLabel();
76 emailTextField = new JTextField(10);
77 phoneLabel = new JLabel();
78 phoneTextField = new JTextField(10);
79 queryPanel = new JPanel();
80 queryLabel = new JLabel();
81 queryTextField = new JTextField(10);
82 queryButton = new JButton();
83 browseButton = new JButton();
84 insertButton = new JButton();
85
86 setLayout(new FlowLayout(FlowLayout.CENTER, 10, 10));
87 setSize(400, 300);
```

Fig. 28.32 | A simple address book. (Part 2 of 8.)

```
88 setResizable(false);
89
90 navigatePanel.setLayout(
91 new BoxLayout(navigatePanel, BoxLayout.X_AXIS));
92
93 previousButton.setText("Previous");
94 previousButton.setEnabled(false);
95 previousButton.addActionListener(
96 new ActionListener()
97 {
98 public void actionPerformed(ActionEvent evt)
99 {
100 previousButtonActionPerformed(evt);
101 } // end method actionPerformed
102 } // end anonymous inner class
103); // end call to addActionListener
104
105 navigatePanel.add(previousButton);
106 navigatePanel.add(Box.createHorizontalStrut(10));
107
108 indexTextField.setHorizontalAlignment(
109 JTextField.CENTER);
110 indexTextField.addActionListener(
111 new ActionListener()
112 {
113 public void actionPerformed(ActionEvent evt)
114 {
115 indexTextFieldActionPerformed(evt);
116 } // end method actionPerformed
117 } // end anonymous inner class
118); // end call to addActionListener
119
120 navigatePanel.add(indexTextField);
121 navigatePanel.add(Box.createHorizontalStrut(10));
122
123 ofLabel.setText("of");
124 navigatePanel.add(ofLabel);
125 navigatePanel.add(Box.createHorizontalStrut(10));
126
127 maxTextField.setHorizontalAlignment(
128 JTextField.CENTER);
129 maxTextField.setEditable(false);
130 navigatePanel.add(maxTextField);
131 navigatePanel.add(Box.createHorizontalStrut(10));
132
133 nextButton.setText("Next");
134 nextButton.setEnabled(false);
135 nextButton.addActionListener(
136 new ActionListener()
137 {
138 public void actionPerformed(ActionEvent evt)
139 {
140 nextButtonActionPerformed(evt);
141 } // end method actionPerformed
142 } // end call to addActionListener
143); // end call to add
144
145 scrollPane.setViewportView(navigatePanel);
146
147 return scrollPane;
148}
```

---

**Fig. 28.32** | A simple address book. (Part 3 of 8.)

```
142 } // end anonymous inner class
143); // end call to addActionListener
144
145 navigatePanel.add(nextButton);
146 add(navigatePanel);
147
148 displayPanel.setLayout(new GridLayout(5, 2, 4, 4));
149
150 idLabel.setText("Address ID:");
151 displayPanel.add(idLabel);
152
153 idTextField.setEditable(false);
154 displayPanel.add(idTextField);
155
156 firstNameLabel.setText("First Name:");
157 displayPanel.add(firstNameLabel);
158 displayPanel.add(firstNameTextField);
159
160 lastNameLabel.setText("Last Name:");
161 displayPanel.add(lastNameLabel);
162 displayPanel.add(lastNameTextField);
163
164 emailLabel.setText("Email:");
165 displayPanel.add(emailLabel);
166 displayPanel.add(emailTextField);
167
168 phoneLabel.setText("Phone Number:");
169 displayPanel.add(phoneLabel);
170 displayPanel.add(phoneTextField);
171 add(displayPanel);
172
173 queryPanel.setLayout(
174 new BoxLayout(queryPanel, BoxLayout.X_AXIS));
175
176 queryPanel.setBorder(BorderFactory.createTitledBorder(
177 "Find an entry by last name"));
178 queryLabel.setText("Last Name:");
179 queryPanel.add(Box.createHorizontalStrut(5));
180 queryPanel.add(queryLabel);
181 queryPanel.add(Box.createHorizontalStrut(10));
182 queryPanel.add(queryTextField);
183 queryPanel.add(Box.createHorizontalStrut(10));
184
185 queryButton.setText("Find");
186 queryButton.addActionListener(
187 new ActionListener()
188 {
189 public void actionPerformed(ActionEvent evt)
190 {
191 queryButtonActionPerformed(evt);
192 } // end method actionPerformed
193 } // end anonymous inner class
194); // end call to addActionListener
195
```

Fig. 28.32 | A simple address book. (Part 4 of 8.)

```
196 queryPanel.add(queryButton);
197 queryPanel.add(Box.createHorizontalStrut(5));
198 add(queryPanel);
199
200 browseButton.setText("Browse All Entries");
201 browseButton.addActionListener(
202 new ActionListener()
203 {
204 public void actionPerformed(ActionEvent evt)
205 {
206 browseButtonActionPerformed(evt);
207 } // end method actionPerformed
208 } // end anonymous inner class
209); // end call to addActionListener
210
211 add(browseButton);
212
213 insertButton.setText("Insert New Entry");
214 insertButton.addActionListener(
215 new ActionListener()
216 {
217 public void actionPerformed(ActionEvent evt)
218 {
219 insertButtonActionPerformed(evt);
220 } // end method actionPerformed
221 } // end anonymous inner class
222); // end call to addActionListener
223
224 add(insertButton);
225
226 addWindowListener(
227 new WindowAdapter()
228 {
229 public void windowClosing(WindowEvent evt)
230 {
231 personQueries.close(); // close database connection
232 System.exit(0);
233 } // end method windowClosing
234 } // end anonymous inner class
235); // end call to addWindowListener
236
237 setVisible(true);
238 } // end no-argument constructor
239
240 // handles call when previousButton is clicked
241 private void previousButtonActionPerformed(ActionEvent evt)
242 {
243 currentEntryIndex--;
244
245 if (currentEntryIndex < 0)
246 currentEntryIndex = numberEntries - 1;
247
248 indexTextField.setText("" + (currentEntryIndex + 1));
```

---

Fig. 28.32 | A simple address book. (Part 5 of 8.)

```
249 indexTextFieldActionPerformed(evt);
250 } // end method previousButtonActionPerformed
251
252 // handles call when nextButton is clicked
253 private void nextButtonActionPerformed(ActionEvent evt)
254 {
255 currentEntryIndex++;
256
257 if (currentEntryIndex >= numberofEntries)
258 currentEntryIndex = 0;
259
260 indexTextField.setText("" + (currentEntryIndex + 1));
261 indexTextFieldActionPerformed(evt);
262 } // end method nextButtonActionPerformed
263
264 // handles call when queryButton is clicked
265 private void queryButtonActionPerformed(ActionEvent evt)
266 {
267 results =
268 personQueries.getPeopleByLastName(queryTextField.getText());
269 numberofEntries = results.size();
270
271 if (numberofEntries != 0)
272 {
273 currentEntryIndex = 0;
274 currentEntry = results.get(currentEntryIndex);
275 idTextField.setText("" + currentEntry.getAddressID());
276 firstNameTextField.setText(currentEntry.getFirstName());
277 lastNameTextField.setText(currentEntry.getLastName());
278 emailTextField.setText(currentEntry.getEmail());
279 phoneTextField.setText(currentEntry.getPhoneNumber());
280 maxTextField.setText("" + numberofEntries);
281 indexTextField.setText("" + (currentEntryIndex + 1));
282 nextButton.setEnabled(true);
283 previousButton.setEnabled(true);
284 } // end if
285 else
286 browseButtonActionPerformed(evt);
287 } // end method queryButtonActionPerformed
288
289 // handles call when a new value is entered in indexTextField
290 private void indexTextFieldActionPerformed(ActionEvent evt)
291 {
292 currentEntryIndex =
293 (Integer.parseInt(indexTextField.getText()) - 1);
294
295 if (numberofEntries != 0 && currentEntryIndex < numberofEntries)
296 {
297 currentEntry = results.get(currentEntryIndex);
298 idTextField.setText("" + currentEntry.getAddressID());
299 firstNameTextField.setText(currentEntry.getFirstName());
300 lastNameTextField.setText(currentEntry.getLastName());
301 emailTextField.setText(currentEntry.getEmail());
302 phoneTextField.setText(currentEntry.getPhoneNumber());
```

Fig. 28.32 | A simple address book. (Part 6 of 8.)

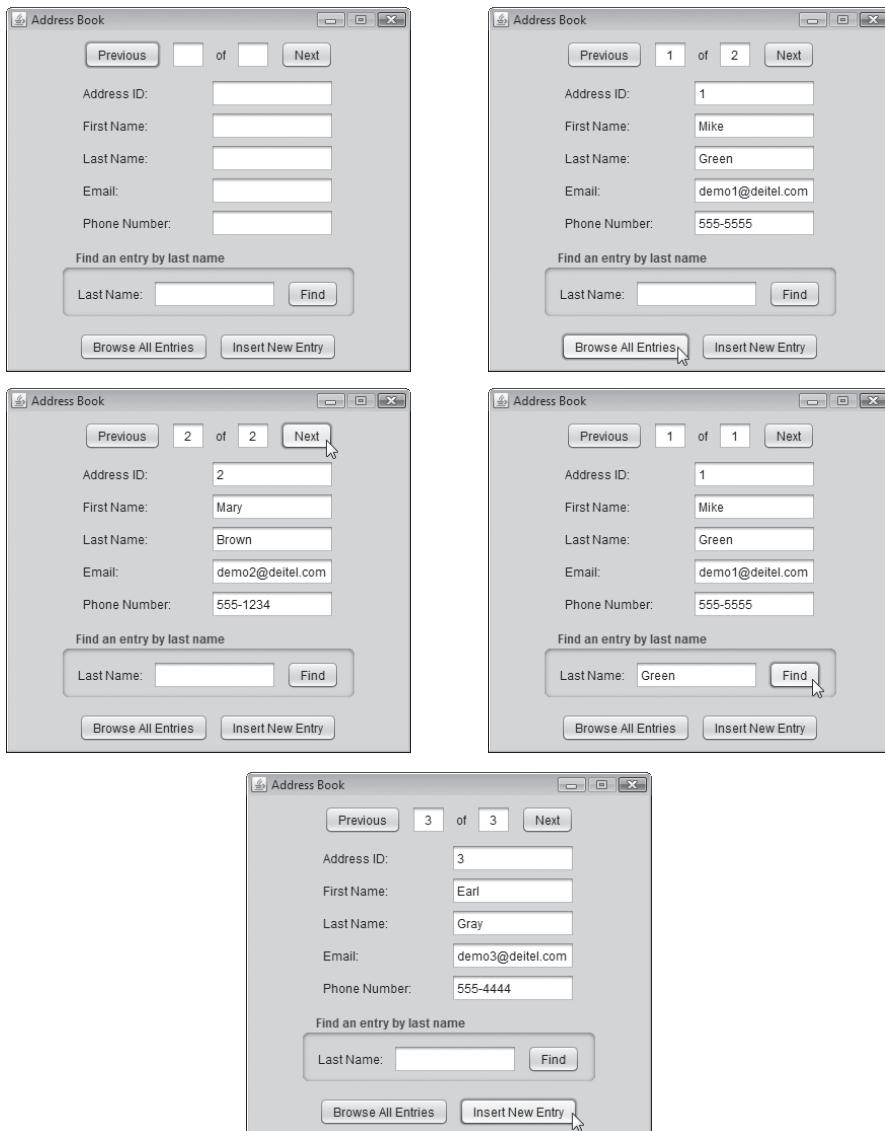
```
303 maxTextField.setText("" + numberOfEntries);
304 indexTextField.setText("" + (currentEntryIndex + 1));
305 } // end if
306 } // end method indexTextFieldActionPerformed
307
308 // handles call when browseButton is clicked
309 private void browseButtonActionPerformed(ActionEvent evt)
310 {
311 try
312 {
313 results = personQueries.getAllPeople();
314 numberOfEntries = results.size();
315
316 if (numberOfEntries != 0)
317 {
318 currentEntryIndex = 0;
319 currentEntry = results.get(currentEntryIndex);
320 idTextField.setText("" + currentEntry.getAddressID());
321 firstNameTextField.setText(currentEntry.getFirstName());
322 lastNameTextField.setText(currentEntry.getLastName());
323 emailTextField.setText(currentEntry.getEmail());
324 phoneTextField.setText(currentEntry.getPhoneNumber());
325 maxTextField.setText("" + numberOfEntries);
326 indexTextField.setText("" + (currentEntryIndex + 1));
327 nextButton.setEnabled(true);
328 previousButton.setEnabled(true);
329 } // end if
330 } // end try
331 catch (Exception e)
332 {
333 e.printStackTrace();
334 } // end catch
335 } // end method browseButtonActionPerformed
336
337 // handles call when insertButton is clicked
338 private void insertButtonActionPerformed(ActionEvent evt)
339 {
340 int result = personQueries.addPerson(firstNameTextField.getText(),
341 lastNameTextField.getText(),
342 emailTextField.getText(),
343 phoneTextField.getText());
344
345 if (result == 1)
346 JOptionPane.showMessageDialog(this, "Person added!",
347 "Person added", JOptionPane.PLAIN_MESSAGE);
348 else
349 JOptionPane.showMessageDialog(this, "Person not added!",
350 "Error", JOptionPane.PLAIN_MESSAGE);
351
352 browseButtonActionPerformed(evt);
353 } // end method insertButtonActionPerformed
354
355 // main method
356 public static void main(String args[])
357 {
```

Fig. 28.32 | A simple address book. (Part 7 of 8.)

```

357 new AddressBookDisplay();
358 } // end method main
359 } // end class AddressBookDisplay

```



**Fig. 28.32** | A simple address book. (Part 8 of 8.)

To add a new entry into the `AddressBook` database, the user can enter the first name, last name, email and phone number (the `AddressID` will autoincrement) in the `JTextFields` and press the `Insert New Entry JButton`. When the user presses `Insert New Entry`, the `insertButtonActionPerformed` handler (lines 338–352) is called. Lines 340–342 call the method `addPerson` on the `PersonQueries` object to add a new entry to the database. Line

351 calls `browseButtonActionPerformed` to obtain the updated set of people in the address book and update the GUI accordingly.

The user can then view different entries by pressing the `Previous JButton` or `Next JButton`, which results in calls to methods `previousButtonActionPerformed` (lines 241–250) or `nextButtonActionPerformed` (lines 253–262), respectively. Alternatively, the user can enter a number in the `indexTextField` and press *Enter* to view a particular entry. This results in a call to method `indexTextFieldActionPerformed` (lines 290–306) to display the specified record.

## 28.12 Stored Procedures

Many database management systems can store individual SQL statements or sets of SQL statements in a database, so that programs accessing that database can invoke them. Such named collections of SQL statements are called **stored procedures**. JDBC enables programs to invoke stored procedures using objects that implement the interface **CallableStatement**. `CallableStatements` can receive arguments specified with the methods inherited from interface `PreparedStatement`. In addition, `CallableStatements` can specify **output parameters** in which a stored procedure can place return values. Interface `CallableStatement` includes methods to specify which parameters in a stored procedure are output parameters. The interface also includes methods to obtain the values of output parameters returned from a stored procedure. To learn more about `CallableStatements`, visit

```
java.sun.com/javase/6/docs/technotes/guides/jdbc/getstart/
callablestatement.html#999652
```



### Portability Tip 28.6

*Although the syntax for creating stored procedures differs across database management systems, the interface CallableStatement provides a uniform interface for specifying input and output parameters for stored procedures and for invoking stored procedures.*



### Portability Tip 28.7

*According to the Java API documentation for interface CallableStatement, for maximum portability between database systems, programs should process the update counts (which indicate how many rows were updated) or ResultSets returned from a CallableStatement before obtaining the values of any output parameters.*

## 28.13 Transaction Processing

Many database applications require guarantees that a series of database insertions, updates and deletions executes properly before the applications continue processing the next database operation. For example, when you transfer money electronically between bank accounts, several factors determine if the transaction is successful. You begin by specifying the source account and the amount you wish to transfer from that account to a destination account. Next, you specify the destination account. The bank checks the source account to determine whether its funds are sufficient to complete the transfer. If so, the bank withdraws the specified amount and, if all goes well, deposits it into the destination account to complete the transfer. What happens if the transfer fails after the bank withdraws the money from the source account? In a proper banking system, the bank redeposits the money

in the source account. How would you feel if the money was subtracted from your source account and the bank *did not* deposit the money in the destination account?

**Transaction processing** enables a program that interacts with a database to treat a database operation (or set of operations) as a single operation. Such an operation also is known as an **atomic operation** or a **transaction**. At the end of a transaction, a decision can be made either to **commit the transaction** or **roll back the transaction**. Committing the transaction finalizes the database operation(s); all insertions, updates and deletions performed as part of the transaction cannot be reversed without performing a new database operation. Rolling back the transaction leaves the database in its state prior to the database operation. This is useful when a portion of a transaction fails to complete properly. In our bank-account-transfer discussion, the transaction would be rolled back if the deposit could not be made into the destination account.

Java provides transaction processing via methods of interface `Connection`. Method `setAutoCommit` specifies whether each SQL statement commits after it completes (a `true` argument) or whether several SQL statements should be grouped as a transaction (a `false` argument). If the argument to `setAutoCommit` is `false`, the program must follow the last SQL statement in the transaction with a call to `Connection` method `commit` (to commit the changes to the database) or `Connection` method `rollback` (to return the database to its state prior to the transaction). Interface `Connection` also provides method `getAutoCommit` to determine the autocommit state for the `Connection`.

## 28.14 Wrap-Up

In this chapter, you learned basic database concepts, how to interact with data in a database using SQL and how to use JDBC to allow Java applications to manipulate MySQL and Java DB databases. You learned about the SQL commands `SELECT`, `INSERT`, `UPDATE` and `DELETE`, as well as clauses such as `WHERE`, `ORDER BY` and `INNER JOIN`. You learned the explicit steps for obtaining a `Connection` to the database, creating a `Statement` to interact with the database's data, executing the statement and processing the results. Then you used a `ResultSet` to simplify the process of connecting to a database and creating statements. You used `PreparedStatement`s to create precompiled SQL statements. You also learned how to create and configure databases in both MySQL and Java DB by using predefined SQL scripts. We also provided overviews of `CallableStatement`s and transaction processing. In the next chapter, you'll learn about web application development with JavaServer Faces.

## 28.15 Web Resources

[java.sun.com/javase/technologies/database/index.jsp](http://java.sun.com/javase/technologies/database/index.jsp)

Sun Microsystems, Java SE database technologies home page.

[java.sun.com/docs/books/tutorial/jdbc/index.html](http://java.sun.com/docs/books/tutorial/jdbc/index.html)

*The Java Tutorial's* JDBC track.

[developers.sun.com/product/jdbc/drivers](http://developers.sun.com/product/jdbc/drivers)

Sun Microsystems search engine for locating JDBC drivers.

[www.sql.org](http://www.sql.org)

This SQL portal provides links to many resources, including SQL syntax, tips, tutorials, books, magazines, discussion groups, companies with SQL services, SQL consultants and free software.

[www.datadirect.com/developer/jdbc/topics/perfoptjdbc/index.ssp](http://www.datadirect.com/developer/jdbc/topics/perfoptjdbc/index.ssp)

White paper that discusses designing a good JDBC application.

[java.sun.com/javase/6/docs/technotes/guides/jdbc/index.html](http://java.sun.com/javase/6/docs/technotes/guides/jdbc/index.html)  
Sun Microsystems JDBC API documentation.

[www.mysql.com](http://www.mysql.com)

This site is the MySQL database home page. You can download the latest versions of MySQL and MySQL Connector/J and access their online documentation.

[dev.mysql.com/doc/refman/5.0/en/index.html](http://dev.mysql.com/doc/refman/5.0/en/index.html)

MySQL reference manual.

[java.sun.com/javase/6/docs/technotes/guides/jdbc/getstart/rowsetImpl.html](http://java.sun.com/javase/6/docs/technotes/guides/jdbc/getstart/rowsetImpl.html)

Overviews the RowSet interface and its subinterfaces. This site also discusses the reference implementations of these interfaces from Sun and their usage.

[java.sun.com/developer/Books/JDBCTutorial/chapter5.html](http://java.sun.com/developer/Books/JDBCTutorial/chapter5.html)

Chapter 5 (RowSet Tutorial) of the book *The JDBC 2.0 API Tutorial and Reference, Second Edition*.

## Summary

### Section 28.1 Introduction

- A database is an integrated collection of data. A database management system (DBMS) provides mechanisms for storing, organizing, retrieving and modifying data for many users.
- Today's most popular database management systems are relational database systems.
- SQL is the international standard language used to query and manipulate relational data.
- Programs connect to, and interact with, relational databases via an interface—software that facilitates communications between a database management system and a program.
- A JDBC driver enables Java applications to connect to a database in a particular DBMS and allows you to retrieve and manipulate database data.

### Section 28.2 Relational Databases

- A relational database stores data in tables. Tables are composed of rows, and rows are composed of columns in which values are stored.
- A primary key provides a unique value that cannot be duplicated in other rows of the same table.
- Each column of a table represents a different attribute.
- The primary key can be composed of more than one column.
- Every column in a primary key must have a value, and the value of the primary key must be unique. This is known as the Rule of Entity Integrity.
- A one-to-many relationship between tables indicates that a row in one table can have many related rows in a separate table.
- A foreign key is a column in a table that must match the primary-key column in another table. This is known as the Rule of Referential Integrity.
- Foreign keys enable information from multiple tables to be joined together. There is a one-to-many relationship between a primary key and its corresponding foreign key.

### Section 28.4.1 Basic SELECT Query

- The basic form of a query is

**SELECT \* FROM tableName**

where the asterisk (\*) indicates that all columns from *tableName* should be selected, and *tableName* specifies the table in the database from which rows will be retrieved.

- To retrieve specific columns, replace the \* with a comma-separated list of column names.

### Section 28.4.2 WHERE Clause

- The optional WHERE clause in a query specifies the selection criteria for the query. The basic form of a query with selection criteria is

```
SELECT columnName1, columnName2, ... FROM tableName WHERE criteria
```

- The WHERE clause can contain operators <, >, <=, >=, =, <> and LIKE. Operator LIKE is used for string pattern matching with wildcard characters percent (%) and underscore (\_).
- A percent character (%) in a pattern indicates that a string matching the pattern can have zero or more characters at the percent character's location in the pattern.
- An underscore (\_) in the pattern string indicates a single character at that position in the pattern.

### Section 28.4.3 ORDER BY Clause

- A query's result can be sorted with the ORDER BY clause. The simplest form of an ORDER BY clause is

```
SELECT columnName1, columnName2, ... FROM tableName ORDER BY column ASC
SELECT columnName1, columnName2, ... FROM tableName ORDER BY column DESC
```

where ASC specifies ascending order, DESC specifies descending order and column specifies the column on which the sort is based. The default sorting order is ascending, so ASC is optional.

- Multiple columns can be used for ordering purposes with an ORDER BY clause of the form

```
ORDER BY column1 sortingOrder, column2 sortingOrder, ...
```

- The WHERE and ORDER BY clauses can be combined in one query. If used, ORDER BY must be the last clause in the query.

### Section 28.4.4 Merging Data from Multiple Tables: INNER JOIN

- An INNER JOIN merges rows from two tables by matching values in columns that are common to the tables. The basic form for the INNER JOIN operator is:

```
SELECT columnName1, columnName2, ...
FROM table1
INNER JOIN table2
ON table1.columnName = table2.columnName
```

The ON clause specifies the columns from each table that are compared to determine which rows are joined. If a SQL statement uses columns with the same name from multiple tables, the column names must be fully qualified by prefixing them with their table names and a dot (.).

### Section 28.4.5 INSERT Statement

- An INSERT statement inserts a new row into a table. The basic form of this statement is

```
INSERT INTO tableName (columnName1, columnName2, ..., columnNameN)
VALUES (value1, value2, ..., valueN)
```

where tableName is the table in which to insert the row. The tableName is followed by a comma-separated list of column names in parentheses. The list of column names is followed by the SQL keyword VALUES and a comma-separated list of values in parentheses.

- SQL uses single quotes ('') to delimit strings. To specify a string containing a single quote in SQL, escape the single quote with another single quote (i.e., ''').

### Section 28.4.6 UPDATE Statement

- An UPDATE statement modifies data in a table. The basic form of an UPDATE statement is

```
UPDATE tableName
SET columnName1 = value1, columnName2 = value2, ..., columnNameN = valueN
WHERE criteria
```

where *tableName* is the table in which to update data. The *tableName* is followed by keyword **SET** and a comma-separated list of *columnName = value* pairs. The optional **WHERE** clause *criteria* determines which rows to update.

#### **Section 28.4.7 DELETE Statement**

- A **DELETE** statement removes rows from a table. The simplest form for a **DELETE** statement is

```
DELETE FROM tableName WHERE criteria
```

where *tableName* is the table from which to delete a row (or rows). The optional **WHERE** *criteria* determines which rows to delete. If this clause is omitted, all the table's rows are deleted.

#### **Section 28.8.1 Connecting to and Querying a Database**

- Package `java.sql` contains classes and interfaces for accessing relational databases in Java.
- A **Connection** object manages the connection between a Java program and a database. Connection objects enable programs to create SQL statements that access data.
- **DriverManager** method `getConnection` attempts to connect to a database at a URL that specifies the protocol for communication, the subprotocol for communication and the database name.
- **Connection** method `createStatement` creates a **Statement** object, which can be used to submit SQL statements to the database.
- **Statement** method `executeQuery` executes a query and returns a **ResultSet** object containing the query result. **ResultSet** methods enable a program to manipulate query results.
- A **ResultSetMetaData** object describes a **ResultSet**'s contents. Programs can use metadata programmatically to obtain information about the **ResultSet** column names and types.
- **ResultSetMetaData** method `getColumnCount` retrieves the number of **ResultSet** columns.
- **ResultSet** method `next` positions the **ResultSet** cursor to the next row and returns `true` if the row exists; otherwise, it returns `false`. This method must be called to begin processing a **ResultSet** because the cursor is initially positioned before the first row.
- It's possible to extract each **ResultSet** column as a specific Java type. **ResultSetMetaData** method `getColumnType` returns a **Types** constant (package `java.sql`) indicating the column's type.
- **ResultSet** `get` methods typically receive as an argument either a column number (as an `int`) or a column name (as a `String`) indicating which column's value to obtain.
- **ResultSet** row and column numbers start at 1.
- Each **Statement** object can open only one **ResultSet** at a time. When a **Statement** returns a new **ResultSet**, the **Statement** closes the prior **ResultSet**.
- **Connection** method `createStatement` has an overloaded version that receives the result type and the result concurrency. The result type specifies whether the **ResultSet**'s cursor is able to scroll in both directions or forward only and whether the **ResultSet** is sensitive to changes. The result concurrency specifies whether the **ResultSet** can be updated with **ResultSet**'s update methods.
- Some JDBC drivers do not support scrollable or updatable **ResultSet**s.

#### **Section 28.8.2 Querying the books Database**

- **TableModel** method `getColumnClass` returns a **Class** object that represents the superclass of all objects in a particular column. A **JTable** uses this information to set up the default cell renderer and cell editor for that column in a **JTable**.
- **ResultSetMetaData** method `getColumnName` obtains a column's fully qualified class name.
- **TableModel** method `getColumnCount` returns the number of columns in the underlying **ResultSet**.
- **TableModel** method `getColumnName` returns the column name in the underlying **ResultSet**.
- **ResultSetMetaData** method `getColumnName` obtains the column name from the **ResultSet**.

- `TableModel` method `getRowCount` returns the number of rows in the model's `ResultSet`.
- `TableModel` method `getValueAt` returns the `Object` at a particular row and column of the model's underlying `ResultSet`.
- `ResultSet` method `absolute` positions the `ResultSet` cursor at a specific row.
- `AbstractTableModel` method `fireTableStructureChanged` notifies any `JTable` using a particular `TableModel` object as its model that the data in the model has changed.

### ***Section 28.9 RowSet Interface***

- Interface `RowSet` configures a database connection and executes a query automatically.
- A connected `RowSet` remains connected to the database while the object is in use. A disconnected `RowSet` connects, executes a query, then closes the connection.
- `JdbcRowSet`, a connected `RowSet`, wraps a `ResultSet` object and allows you to scroll and update its rows. Unlike a `ResultSet` object, a `JdbcRowSet` object is scrollable and updatable by default.
- `CachedRowSet`, a disconnected `RowSet`, caches a `ResultSet`'s data in memory. A `CachedRowSet` is scrollable and updatable. A `CachedRowSet` is also serializable.

### ***Section 28.10 Java DB/Apache Derby***

- As of JDK 6, Sun Microsystems bundles the open-source, pure Java database Java DB (the Sun branded version of Apache Derby) with the JDK.
- Java DB has both an embedded version and a network version.

### ***Section 28.11 PreparedStatements***

- `PreparedStatement`s are compiled, so they execute more efficiently than `Statement`s.
- `PreparedStatement`s can have parameters, so the same query can execute with different arguments.
- A parameter is specified with a question mark (?) in the SQL statement. Before executing a `PreparedStatement`, you must use `PreparedStatement`'s `set` methods to specify the arguments.
- `PreparedStatement` method `setString`'s first argument represents the parameter number being set and the second argument is that parameter's value.
- Parameter numbers are counted from 1, starting with the first question mark (?).
- Method `setString` automatically escapes `String` parameter values as necessary.
- Interface `PreparedStatement` provides `set` methods for each supported SQL type.
- An identity column is the SQL standard way to represent an autoincremented column. The SQL `IDENTITY` keyword marks a column as an identity column.

### ***Section 28.12 Stored Procedures***

- JDBC enables programs to invoke stored procedures using `CallableStatement` objects.
- `CallableStatement` can specify input parameters. `CallableStatement` can specify output parameters in which a stored procedure can place return values.

### ***Section 28.13 Transaction Processing***

- Transaction processing enables a program that interacts with a database to treat a database operation (or set of operations) as a single operation—known as an atomic operation or a transaction.
- At the end of a transaction, a decision can be made to either commit or roll back the transaction.
- Committing a transaction finalizes the database operation(s)—inserts, updates and deletes cannot be reversed without performing a new database operation.
- Rolling back a transaction leaves the database in its state prior to the database operation.
- Java provides transaction processing via methods of interface `Connection`.

- Method `setAutoCommit` specifies whether each SQL statement commits after it completes (a `true` argument) or whether several SQL statements should be grouped as a transaction.
- When autocommit is disabled, the program must follow the last SQL statement in the transaction with a call to `Connection` method `commit` (to commit the changes to the database) or `Connection` method `rollback` (to return the database to its state prior to the transaction).
- Method `getAutoCommit` determines the autocommit state for the `Connection`.

## Terminology

|                                                                        |              |                                                                 |          |
|------------------------------------------------------------------------|--------------|-----------------------------------------------------------------|----------|
| _ SQL wildcard character                                               | CXXXV        | getColumnCount method of <code>TableModel</code>                | CXLIX    |
| * SQL wildcard character                                               | CXXXIV       | getColumnName method of <code>ResultSetMetaData</code>          | CLIV     |
| % SQL wildcard character                                               | CXXXV        | getColumnName method of <code>TableModel</code>                 | CXLIX    |
| <code>absolute</code> method of <code>ResultSet</code>                 | CLV          | getColumnType method of <code>ResultSetMetaData</code>          | CXLVIII  |
| <code>AbstractTableModel</code> class                                  | CXLIX        | getConnection method of <code>DriverManager</code>              | CXLVI    |
| <code>addTableModelListener</code> method of <code>TableModel</code>   | CXLIX        | getInt method of <code>ResultSet</code>                         | CXLVIII  |
| atomic operation                                                       | CLXXXI       | getObject method of <code>ResultSet</code>                      | CXLVIII  |
| autoincremented value                                                  | CXXX         | getRow method of <code>ResultSet</code>                         | CLV      |
| automatic driver discovery                                             | CXLVI        | getRowCount method of <code>TableModel</code>                   | CXLIX    |
| <code>CachedRowSet</code> interface                                    | CLXI         | getValueAt method of <code>TableModel</code>                    | CXLIX    |
| <code>CallableStatement</code> interface                               | CLXXX        | IDENTITY keyword (SQL)                                          | CLXVI    |
| <code>close</code> method of <code>JdbcRowSet</code>                   | CLXIII       | INNER JOIN SQL clause                                           | CXXXVIII |
| column in a database table                                             | CXXIX        | INSERT SQL statement                                            | CXXXIX   |
| <code>com.sun.rowset</code> package                                    | CLXIII       | Java DB                                                         | CLXIII   |
| <code>commit</code> a transaction                                      | CLXXXI       | <code>javax.sql.rowset</code> package                           | CLXI     |
| <code>commit</code> method of interface <code>Connection</code>        | CLXXXI       | <code>javax.swing.table</code> package                          | CLX      |
| <code>connected RowSet</code>                                          | CLXI         | JDBC API                                                        | CXXVIII  |
| <code>Connection</code> interface                                      | CXLVI        | JDBC driver                                                     | CXXVIII  |
| <code>createStatement</code> method of <code>Connection</code>         | CXLVII       | <code>JdbcRowSet</code> interface                               | CLXI     |
| database                                                               | CXXVIII      | <code>JdbcRowSetImpl</code> class                               | CLXIII   |
| database management system (DBMS)                                      | CXXVIII      | joining database tables                                         | CXXXI    |
| <code>DELETE</code> SQL statement                                      | CXLI         | <code>JTable</code> class                                       | CXLIX    |
| disconnected RowSet                                                    | CLXI         | <code>last</code> method of <code>ResultSet</code>              | CLV      |
| <code>DriverManager</code> class                                       | CXLVI        | LIKE operator (SQL)                                             | CXXXV    |
| entity-relationship (ER) diagram                                       | CXXXII       | metadata                                                        | CXLVII   |
| <code>execute</code> method of <code>JdbcRowSet</code>                 | CLXIII       | MySQL 5.0 Community Edition                                     | CXLII    |
| <code>executeQuery</code> method of interface                          |              | MySQL Connector/J                                               | CXLII    |
| <code>PreparedStatement</code>                                         | CLXXI        | <code>next</code> method of <code>ResultSet</code>              | CXLVIII  |
| <code>executeQuery</code> method of <code>Statement</code>             | CXLVII       | ON clause                                                       | CXXXVIII |
| <code>executeUpdate</code> method of interface                         |              | one-to-many relationship                                        | CXXXIII  |
| <code>PreparedStatement</code>                                         | CLXXII       | ORDER BY SQL clause                                             | CXXXVI   |
| <code>fireTableStructureChanged</code> method of                       |              | output parameter                                                | CLXXX    |
| <code>AbstractTableModel</code>                                        | CLV          | pattern matching                                                | CXXXV    |
| foreign key                                                            | CXXXI        | predicate                                                       | CXXXV    |
| <code>getAutoCommit</code> method of interface <code>Connection</code> |              | <code>PreparedStatement</code> interface                        | CLXV     |
| CLXXXI                                                                 |              | <code>prepareStatement</code> method of interface               |          |
| <code>getColumnClass</code> method of <code>TableModel</code>          | CXLIX        | <code>Connection</code>                                         | CLXVIII  |
| <code>getColumnClassName</code> method of                              |              | primary key                                                     | CXXIX    |
| <code>ResultSetMetaData</code>                                         | CLIV         | qualified name                                                  | CXXXIX   |
| <code>getRowCount</code> method of <code>ResultSetMetaData</code>      | CXLVII, CLIV | query                                                           | CXXVIII  |
|                                                                        |              | <code>regexFilter</code> method of class <code>RowFilter</code> | CLXI     |

|                                                          |                                                         |
|----------------------------------------------------------|---------------------------------------------------------|
| relational database CXXIX                                | setRowFilter method of class JTable CLXI                |
| relational database management system<br>(RDBMS) CXXVIII | setRowSorter method of class JTable CLX                 |
| removeTableModelListener method of<br>TableModel CXLIX   | setString method of interface<br>PreparedStatement CLXV |
| result set concurrency CLIV                              | setUrl method of JdbcRowSet interface CLXIII            |
| result set type CLIII                                    | setUsername method of JdbcRowSet interface<br>CLXIII    |
| ResultSet interface CXLVII                               | SQLException class CXLVII                               |
| ResultSetMetaData interface CXLVII                       | Statement interface CXLVII                              |
| roll back a transaction CLXXXI                           | stored procedure CLXXX                                  |
| rollback method of interface Connection<br>CLXXXI        | Structured Query Language (SQL) CXXVIII                 |
| row in a database table CXXIX                            | subprotocol CXLVII                                      |
| RowFilter class CLXI                                     | table in a database CXXIX                               |
| RowSet interface CLXI                                    | TableModel interface CXLIX                              |
| Rule of Entity Integrity CXXXII                          | TableRowSorter class CLX                                |
| Rule of Referential Integrity CXXXI                      | transaction CLXXXI                                      |
| SELECT SQL keyword CXXXIV                                | transaction processing CLXXXI                           |
| selection criteria CXXXV                                 | Types class CXLVIII                                     |
| SET SQL clause CXL                                       | UPDATE SQL statement CXL                                |
| setAutoCommit method of interface Connection<br>CLXXXI   | VALUES SQL clause CXXXIX                                |
| setCommand method of JdbcRowSet interface<br>CLXIII      | WHERE SQL clause CXXXV                                  |
| setPassword method of JdbcRowSet interface<br>CLXIII     | WindowAdapter class CLXI                                |
|                                                          | windowClosed method of interface<br>WindowListener CLXI |
|                                                          | WindowListener interface CLXI                           |

## Self-Review Exercise

- 28.1** Fill in the blanks in each of the following statements:
- a) The international standard database language is \_\_\_\_\_.
  - b) A table in a database consists of \_\_\_\_\_ and \_\_\_\_\_.
  - c) Statement objects return SQL query results as \_\_\_\_\_ objects.
  - d) The \_\_\_\_\_ uniquely identifies each row in a table.
  - e) SQL keyword \_\_\_\_\_ is followed by the selection criteria that specify the rows to select in a query.
  - f) SQL keywords \_\_\_\_\_ specify the order in which rows are sorted in a query.
  - g) Merging rows from multiple database tables is called \_\_\_\_\_ the tables.
  - h) A(n) \_\_\_\_\_ is an organized collection of data.
  - i) A(n) \_\_\_\_\_ is a set of columns whose values match another table's primary-key values.
  - j) \_\_\_\_\_ method \_\_\_\_\_ is used to obtain a Connection to a database.
  - k) Interface \_\_\_\_\_ helps manage the connection between a Java program and a database.
  - l) A(n) \_\_\_\_\_ object is used to submit a query to a database.
  - m) Unlike a ResultSet object, \_\_\_\_\_ and \_\_\_\_\_ objects are scrollable and updatable by default.
  - n) \_\_\_\_\_, a disconnected RowSet, caches the data of a ResultSet in memory.

## Answers to Self-Review Exercise

- 28.1** a) SQL. b) rows, columns. c) ResultSet. d) primary key. e) WHERE. f) ORDER BY. g) joining. h) database. i) foreign key. j) DriverManager, getConnection. k) Connection. l) Statement. m) JdbcRowSet, CachedRowSet n) CachedRowSet.

## Exercises

**28.2 (Query Application for the books Database)** Using the techniques shown in this chapter, define a complete query application for the books database. Provide the following predefined queries:

- Select all authors from the Authors table.
- Select a specific author and list all books for that author. Include each book's title, year and ISBN. Order the information alphabetically by the author's last then first name.
- Select a specific publisher and list all books published by that publisher. Include the title, year and ISBN. Order the information alphabetically by title.
- Provide any other queries you feel are appropriate.

Display a JComboBox with appropriate names for each predefined query. Also allow users to supply their own queries.

**28.3 (Data Manipulation Application for the books Database)** Define a data-manipulation application for the books database. The user should be able to edit existing data and add new data to the database (obeying referential and entity integrity constraints). Allow the user to edit the database in the following ways:

- Add a new author.
- Edit the existing information for an author.
- Add a new title for an author. (Remember that the book must have an entry in the AuthorISBN table.).
- Add a new entry in the AuthorISBN table to link authors with titles.

**28.4 (Employee Database)** In Section 10.7, we introduced an employee-payroll hierarchy to calculate each employee's payroll. In this exercise, we provide a database of employees that corresponds to the employee-payroll hierarchy. (A SQL script to create the employees database is provided with the examples for this chapter.) Write an application that allows the user to:

- Add employees to the employee table.
- Add payroll information to the appropriate table for each new employee. For example, for a salaried employee add the payroll information to the salariedEmployees table.

Figure 28.33 is the entity-relationship diagram for the employees database.

**28.5 (Employee Database Query Application)** Modify Exercise 28.4 to provide a JComboBox and a JTextArea to allow the user to perform a query that is either selected from the JComboBox or defined in the JTextArea. Sample predefined queries are:

- Select all employees working in Department SALES.
- Select hourly employees working over 30 hours.
- Select all commission employees in descending order of the commission rate.

**28.6 (Employee Database Data Manipulation Application)** Modify Exercise 28.5 to perform the following tasks:

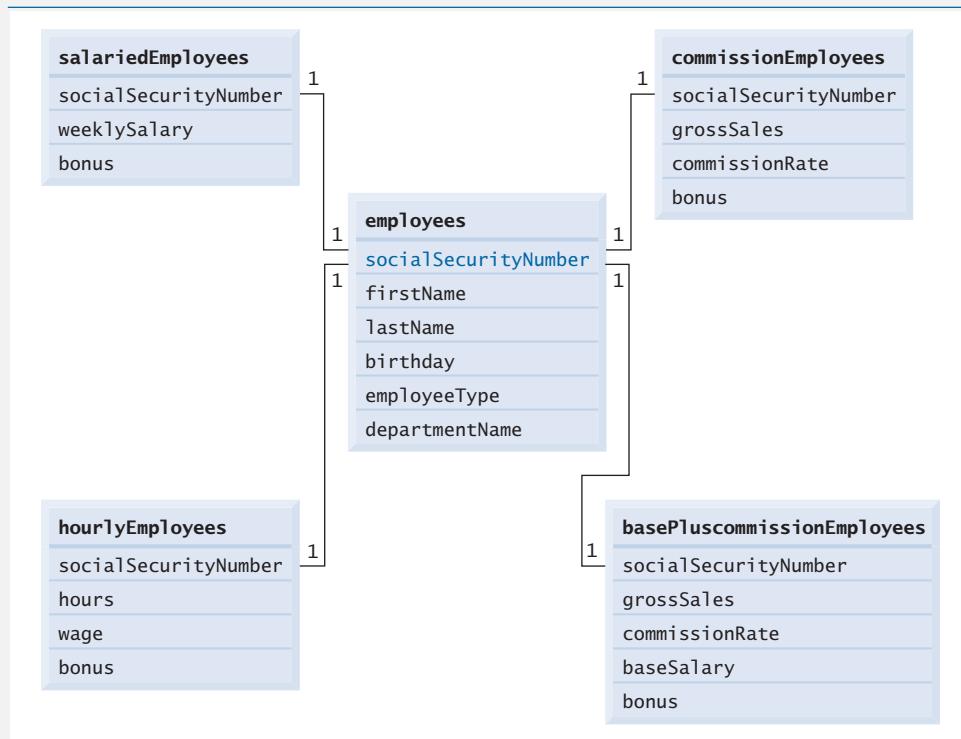
- Increase base salary by 10% for all base-plus-commission employees.
- If the employee's birthday is in the current month, add a \$100 bonus.
- For all commission employees with gross sales over \$10,000, add a \$100 bonus.

**28.7 (AddressBook Modification: Update an Existing Entry)** Modify the program in Figs. 28.30–28.32 to provide a JButton that allows the user to call a method named updatePerson in PersonQueries class to update the current entry in the AddressBook database.

**28.8 (AddressBook Modification: Delete an Existing Entry)** Modify the program of Exercise 28.7 to provide a JButton that allows the user to call a method named deletePerson in PersonQueries class to delete the current entry in the AddressBook database.

**28.9 (ATM Case Study with a Database)** Modify the ATM Case Study (Chapters 12–13) to use an actual database to store the account information. We provide a SQL script to create the BankDa-

tabase, which has a single table consisting of four columns—AccountNumber (an int), PIN (an int), AvailableBalance (a double) and TotalBalance (a double).



**Fig. 28.33** | Table relationships in the employees database.